

System Programming

2. File IO (1): Standard I/O Library

Seung-Ho Lim

Dept. of Computer & Electronic Systems Eng



Linux System Calls

- File descriptor I/O
 - ***open(); close(); creat(), read(); write();***
 - ***seek();*** // random access
 - ***fcntl();*** // for file/record locking
- Process control
- Thread programming
- IPC
- Signal handling
- Memory management
- *Synchronization*
- *Time management*
- *Network socket API (TCP, UDP)*



System Calls & Library Calls for File I/O

- System Calls for File descriptor I/O
 - ***open(); close(); creat(), read(); write();***
 - ***seek();*** // random access
 - ***fcntl();*** // for file/record locking
- Library Calls for File I/O
 - ***fopen(); freopen(); fclose(); fread(); fwrite();***
 - ***fgetc(), fgetchar(); fputc, putchar(); ...***
 - ***fseek(), fprintf(); fscanf();..***



System Calls vs. Library Calls

■ System Calls

- they are entry points into kernel code where their functions are implemented.
- documented in section 2 of the linux manual (e.g. `write(2)` or `man 2 write`)


■ Library Calls

- they are transfers to user code which performs the desired functions.
- documented in section 3 of the linux manual (e.g. `printf(3)`).
- also called *API*(application programming interface)



Library (1)

- A set of compiled object functions for reuse
 - e.g. Graphic Lib., Mathematical Lib., etc.
 - In Linux, generally located in “/lib” or in “/usr/lib”.
 - Only necessary functions(objects) will be **linked** to the user program
- Compile & Linking (review)



Library (2)

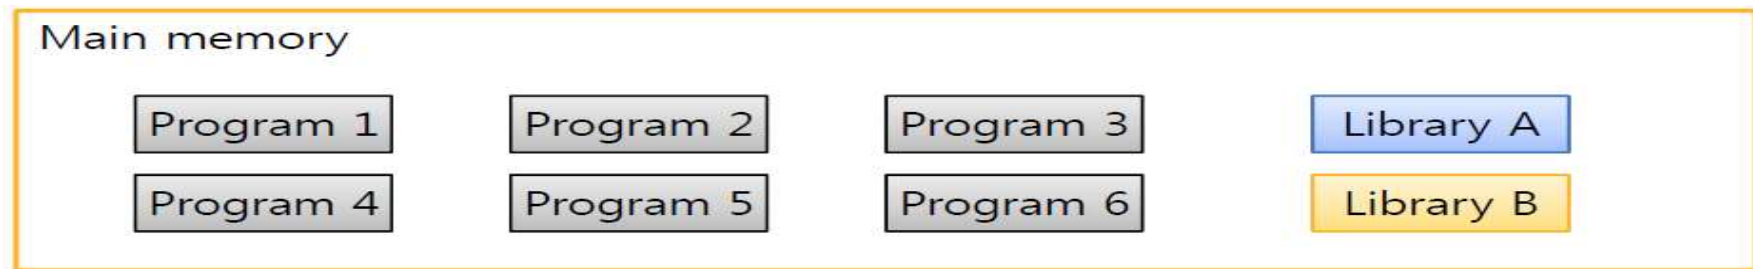
■ Types of libraries.

- Shared library (*.so, *.dll)
 - Only one copy of the function resides in the memory. The function will be shared between several processes. (memory saving)
 - The address of the function will be resolved at run-time. (called **dynamic linking or binding**)
 - A *symbol table* for the dynamic linking exists in memory. (memory overhead).
 - Useful for **server systems**
- Static library (*.a)
 - Necessary functions are added(linked) to each binary program.
 - So, several same copies of a function reside in memory. (overhead)
 - Useful for **embedded systems**

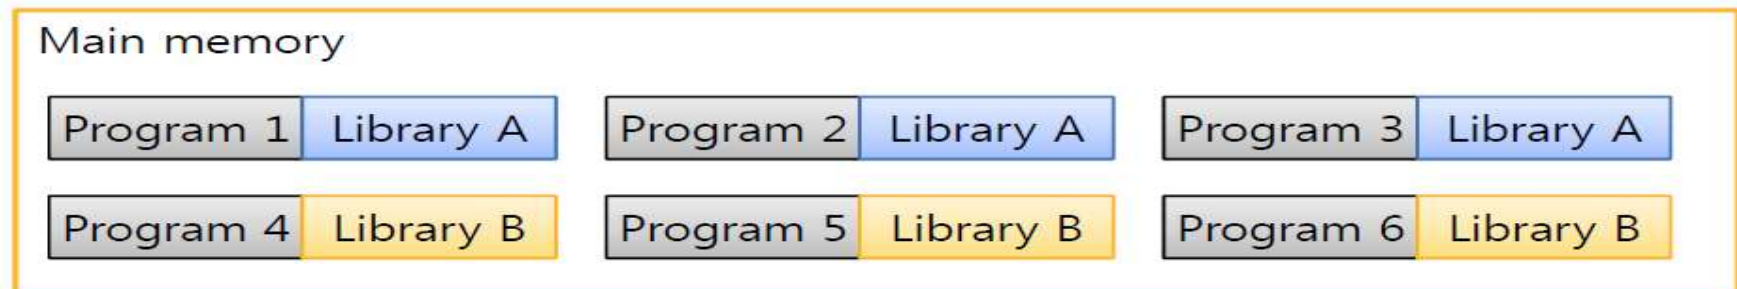


Library (3)

- Executable using **shared** library

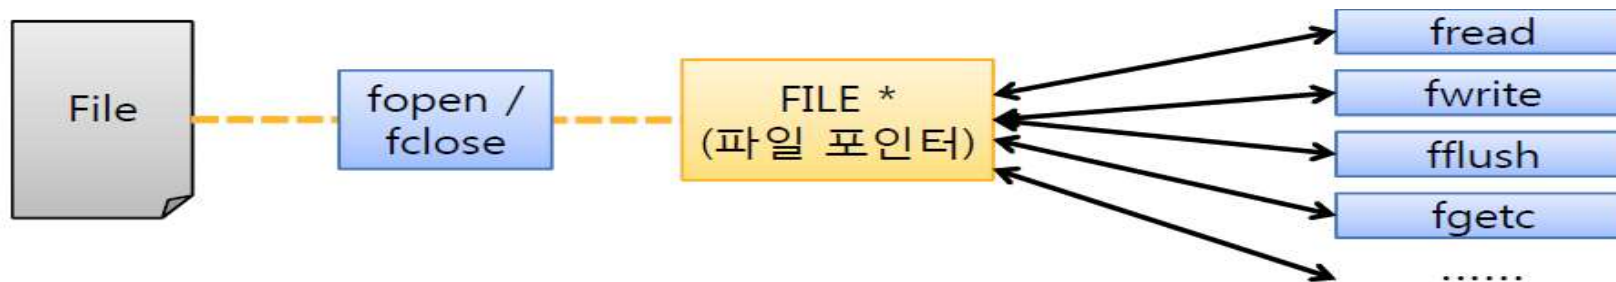


- Executable using **static** library



Standard I/O Library

- `<stdio.h>`
 - a header file which defines symbols and APIs of the standard I/O library (usually for console and files)
- File I/O with the standard I/O library



- I/O devices are mapped to special files
 - Console terminal: ***stdin, stdout, stderr***
 - Console files will be automatically open at run-time.



FILE object in C

- I/O stream object created by standard I/O library
 - accessed by a pointer **FILE***
 - the file stream pointer is used to designate an open file.
 - a file pointer has several system information of an open file.
- `stdin`, `stdout`, `stderr`
 - file stream pointers for the three instances of a console
 - already be opened by the “**shell**” and they are inherited to a user program.



File descriptor

- OS system calls for I/O
 - use file descriptors (NOT FILE*)
 - a file descriptor for an open file is an integer
 - descriptors 0, 1, 2 are assigned to stdin, stdout, stderr
 - for user open files, file descriptors are assigned from 3 in ascending order
 - usually, a user can open 1024 files at maximum
- A standard I/O library function will eventually call the appropriate system call.
 - printf, fprintf, puts, → call **write()**
- Why use standard I/O library?
 - more convenient than simple system calls
 - formatting, library buffering, ...



File stream & File Descriptor

- A file stream is 1:1 mapped to a file descriptor
- Thus, we can get each counterpart information by the following functions

```
#include<stdio.h>  
int fileno(FILE *stream);
```

- returns a file descriptor (number) for the open FILE stream

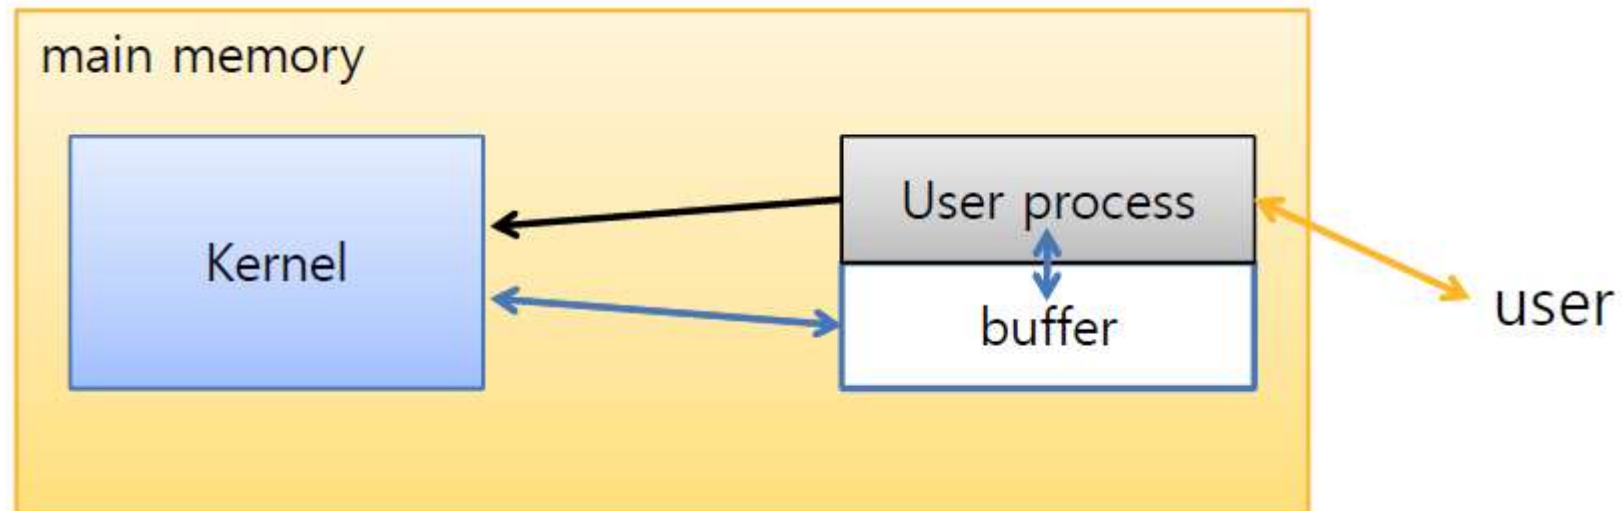
```
#include<stdio.h>  
FILE * fdopen(int fildes, const char *mode);
```

- using the file descriptor of an open file, creates and returns a FILE stream



Library buffering (1)

- Library buffering
 - user-level buffering by library (i.e. user program)
 - reduce the number of system calls
e.g. “DEL” key processing in keyboard input



Library buffering (2)

- Full buffering
 - lib-level buffer for disk blocks (multiple KBs)
 - significantly reduce system calls.
 - For synchronization with the kernel , `fflush()` can be used.
- Line buffering
 - used for console I/O.
 - actual I/O happens when a “`newline`” (enter) appears
 - `getchar()` problem
 - a character is not delivered until entering a “newline”
- Unbuffering
 - no use of library buffer
 - direct delivery to syscalls
 - safe at a power failure.



Library buffering (3)

- Linux library buffering
 - stderr: always **unbuffering**
 - stdin/stdout: always **line buffering**
 - anything else: always **full buffering** (by default)

Set Buffering Type

```
#include <stdio.h>
// set a buffer address that user provides
void setbuf (FILE *stream, char *buf);
    buf : non-NULL address for normal buffering
         NULL if unbuffering
return : none

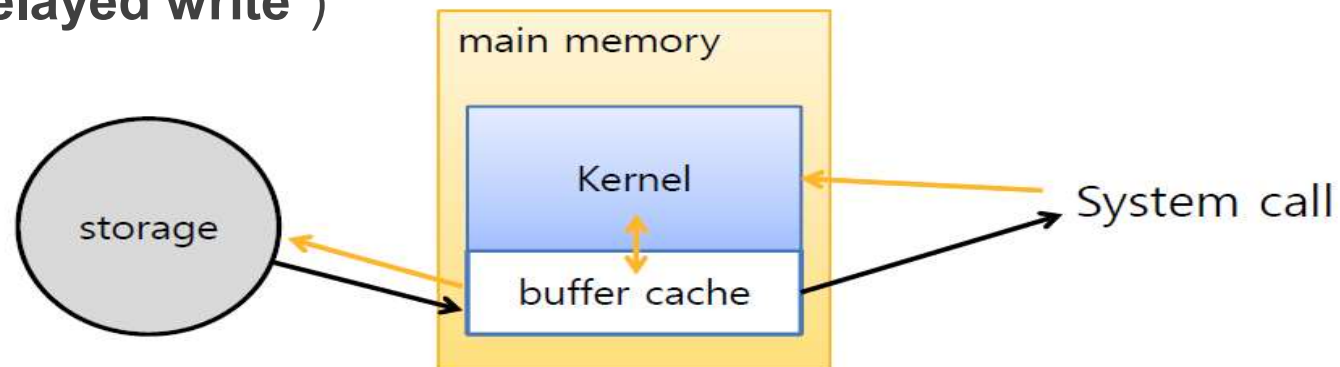
// set a buffer address and buffering type
int setvbuf (FILE *stream, char *buf, int type, size_t size);
    buf : same as the above
    type : the type of buffering
    size : buffer size
return 0 for success, or
      nonzero for an error
```

<i>type</i>	<i>meaning</i>
_IOFBF	Full buffering
_IOLBF	Line buffering
_IONBF	Unbuffering



Kernel Buffering

- Kernel buffering
 - software caching by the kernel.
 - page cache (buffer cache): to reduce disk I/Os.
 - e.g. frequently used disk blocks are kept in the kernel memory (page cache)
 - When reading from a disk
 - try page cache first, if fail do the disk I/O.
 - When writing to a disk
 - write the bytes into the cache, sync to the disk later. (called “**delayed write**”)



fflush

```
#include <stdio.h>
int fflush( FILE *stream);
    return 0 for normal
        EOF for error
```

- flush out the library buffer contents to the kernel. (synchronization)
- due to the buffering, printf (...) does not guarantee the actual output (why?)
- thus, for debugging, write a code as follows

```
printf("something");
fflush(stdout);
```

- for block device I/O (e.g. disk)
 - in block device, a transfer unit b/w disk and kernel is in KBs
 - fflush() moves the contents “lib. buffer” to “page cache”
 - thus, if we want a disk synchronization, use **sync()**
- When a file is closed, fflush() will be done automatically.



I/O buffering & Sync

