

System Programming

3. File IO (2): System Call

Seung-Ho Lim

Dept. of Computer & Electronic Systems Eng



Linux System Calls

- File descriptor I/O
 - ***open(); close(); creat(), read(); write();***
 - ***seek();*** // random access
 - ***fcntl();*** // for file/record locking
- Process control
- Thread programming
- IPC
- Signal handling
- Memory management
- *Synchronization*
- *Time management*
- *Network socket API (TCP, UDP)*



System Calls & Library Calls for File I/O

- System Calls for File descriptor I/O
 - ***open(); close(); creat(), read(); write();***
 - ***seek();*** // random access
 - ***fcntl();*** // for file/record locking
- Library Calls for File I/O
 - ***fopen(); freopen(); fclose(); fread(); fwrite();***
 - ***fgetc(), fgetchar(); fputc, putchar(); ...***
 - ***fseek(), fprintf(); fscanf();..***



System Calls vs. Library Calls

■ System Calls

- they are entry points into kernel code where their functions are implemented.
- documented in section 2 of the linux manual (e.g. `write(2)` or `man 2 write`)

■ Library Calls

- they are transfers to user code which performs the desired functions.
- documented in section 3 of the linux manual (e.g. `printf(3)`).
- also called *API*(application programming interface)



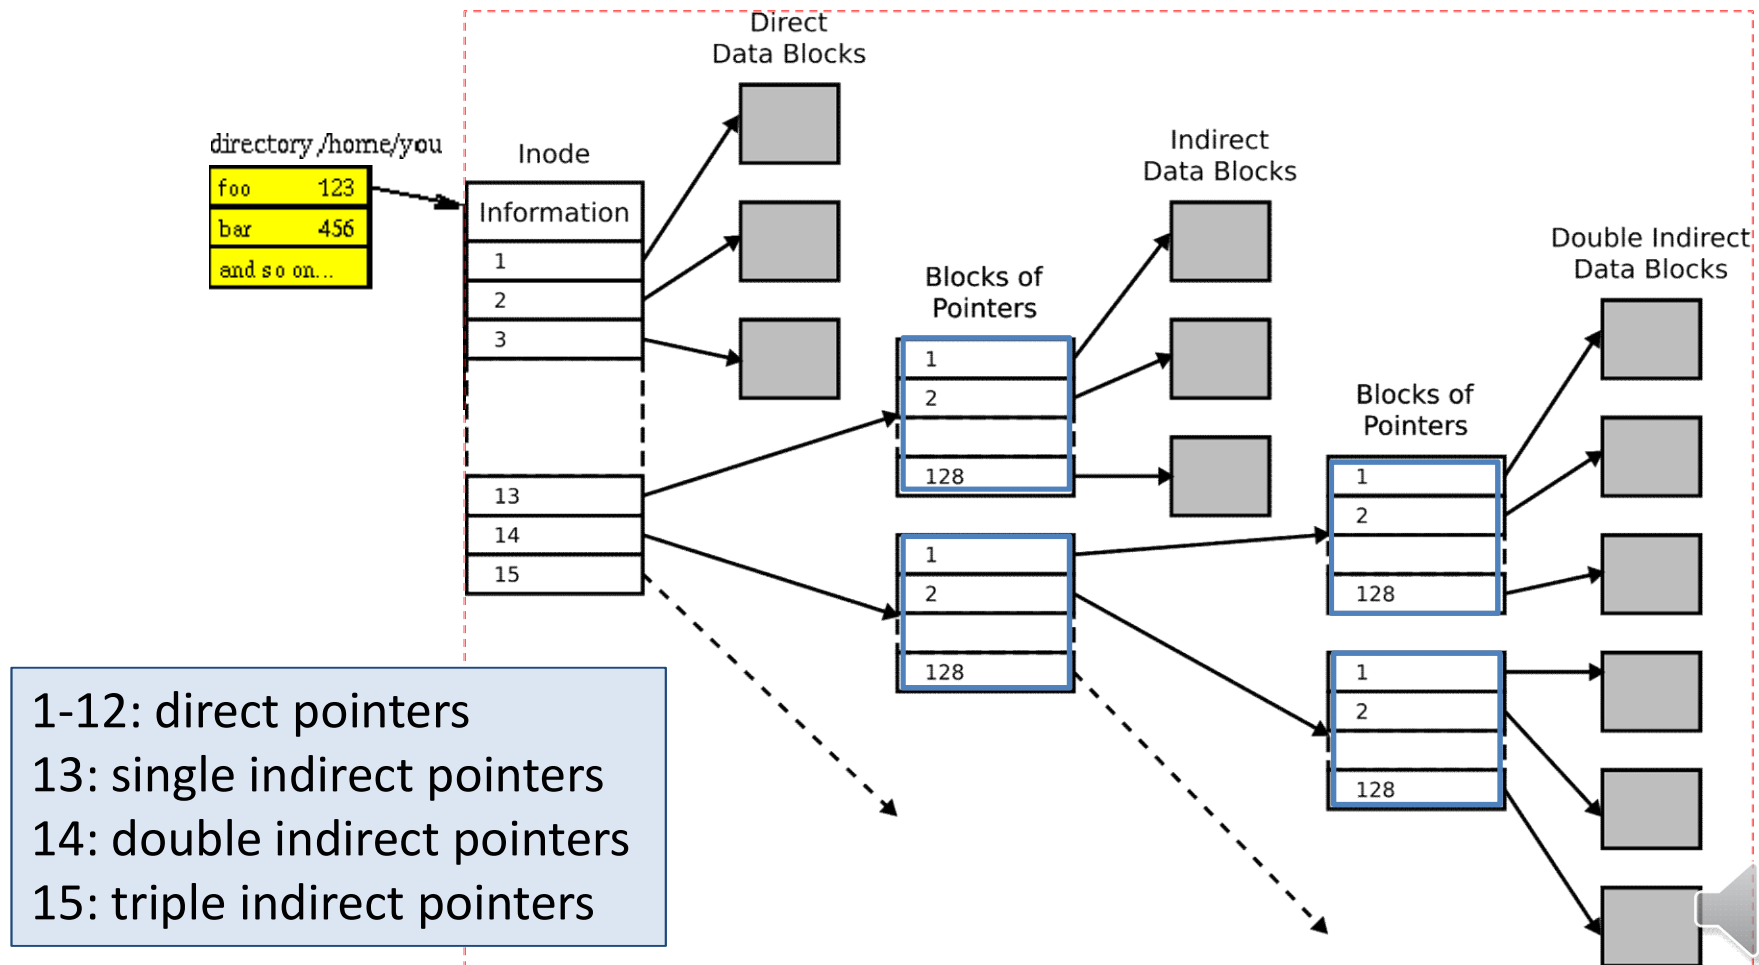
Linux File System (1)

- Each file in a file system has its own *inode*
- An inode is a data structure having all information on a file.
- inodes of all files reside in a *disk*
- inode contents (C struct)
 - file name
 - file type (regular, directory,...)
 - file owner id
 - access permission
rwxr-xr-x (for owner, group, others)
 - creation/modified time
 - file size
 - file data block addr. table (see the next page!)
 - ...



Inode structure example

- block size: 512 byte, block pointer: 4 bytes



Linux File System (2)

■ File types

- regular file
- directory file
- FIFO file (pipe)
- special files (IO devices)
- symbolic link files

A “john” directory file

i	. (john)
j	.. (parent)
k	File name A
l	File name B
m	File name C

■ Directory file

- A directory is just a file whose content is the list of (inode #, file name) in the same directory.
- inode is a data structure which contains all the information about the file and file data blocks
- inode # is a unique file id number in the file system
- “ls -al john” is a shell command that just displays the “john” directory file




Linux File Types

- Ordinary File (Regular File)
 - Text, binary files
- Directory File
 - A file that includes the set of (*file-name*, *inode #*) of the directory.
- Character Special File
 - Character-oriented device (e.g. Keyboard)
- Block Special File
 - Block-oriented device (e.g. HDD file systems, eth0)
- FIFO file
 - Named *pipe* / *Unnamed pipe*
cf. pipe in a process is usually unnamed.
- Symbolic link file
 - a file which points to another file
*cf. **hardlink** is NOT a file.*



File Descriptor (1)

- A *file descriptor* (or *file handle*) is a small, non-negative integer which identifies a file to the kernel.
 - Traditionally, `stdin`, `stdout` and `stderr` are 0, 1 and 2 respectively.



- Relying on “magic numbers” is BAD.
 - Use `STDIN_FILENO`, `STDOUT_FILENO` and `STDERR_FILENO` defined in `unistd.h` or `stdin`, `stdout`, and `stderr` defined in `stdio.h`.



File Descriptor (2)

- Maximum number of files
 - a process can open 1024 files
 - we can check the system resource configuration

```
$ ulimit -a
core file size      (blocks, -c) 0
data seg size       (kbytes, -d) unlimited
scheduling priority (-e) 0
file size           (blocks, -f) unlimited
pending signals     (-i) 194273
max locked memory   (kbytes, -l) 64
max memory size     (kbytes, -m) unlimited
open files          (-n) 1024
pipe size           (512 bytes, -p) 8
.....
```



Basic File I/Os

- 5 fundamental Unix/Linux file I/Os
 - `open(2)`
 - `close(2)`
 - `lseek(2)`
 - `read(2)`
 - `write(2)`

File open (1)

```
#include <fcntl.h>
```

```
int open(const char *path, int oflag);
```

```
int open(const char *path, int oflag, mode_t mode);
```

■ parameters

- *path*: name of the file to open or create
- *oflag*: file open options
- *mode*: access permission (at file creation)

■ return

- *file descriptor* if OK
- 1 on error



File open (2)

- *oflag* options

- must be one of these

option1	meaning	<fcntl.h> defined
O_RDONLY	open for reading only	0
O_WRONLY	open for writing only	1
O_RDWR	open for reading & writing	2

- and can be OR'ed with any of these (by "|")

option2	meaning
O_CREAT	create a file if the file does not exist.
O_EXCL	used with O_CREAT, return an error if the file already exists.
O_TRUNC	if the file exists, make it empty.
O_APPEND	write from the end of the file.
O_SYNC	do disk synchronization when does file I/O.



File access modes

mode	meaning
S_ISUID	set-user-id at execution
S_ISGID	set-group-id at execution
S_ISVTX	set sticky bit
S_IRWXU	owner RWX
S_IRUSR	owner R
S_IWUSR	owner W
S_IXUSR	owner X
S_IRWXG	group RWX
S_IRGRP	group R
S_IWGRP	group W
S_IXGRP	group X
S_IRWXO	others RWX
S_IROTH	others R
S_IWOTH	others W
S_IXOTH	others X

File close

```
#include <unistd.h>

int close(int fd);
```

- parameters
 - *fd*: file descriptor
- return
 - 0 if OK
 - -1 on error

File open example

open-ex.c

```
#include <fcntl.h>

int main(int argc, char * argv[])
{
    FILE *fpo; // file pointer
    int fdo;    // file descriptor

    if(argc != 2) {
        perror(argv[0]);
        return 1;
    }
    if((fdo = open(argv[1], O_RDWR | O_CREAT | O_TRUNC,
                    S_IRUSR | S_IWUSR)) == -1) {
        perror(argv[1]);
        return 1;
    }
    if((fpo = fdopen(fdo, "r+")) == NULL) {
        perror("fdopen");
        return 2;
    }
    fprintf(fpo, "Hello, world! \n");
    fclose(fpo);
}
```

```
$/a.out test.txt
$ cat test.txt
Hello, world!
$
```



File creation

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int creat(const char *path, mode_t mode);
```

- parameters
 - *path* : file path name
 - *mode* : access permission
- return
 - *file descriptor* if OK
 - -1 on error



File seeking

```
#include <sys/types.h>
#include <unistd.h>

off_t lseek(int fd, off_t offset, int whence);
```

■ parameters

- *fd* : file descriptor
- *offset* : offset from the beginning to seek (move)
- *whence* : SEEK_SET, SEEK_CUR, SEEK_END

■ return

- new offset value if OK
- -1 on error



File reading

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t nbyte);
```

■ parameters

- *fd* : file descriptor
- *buf* : buffer address
- *nbyte* : number of bytes to read

■ return

- number of bytes read successfully if OK
- -1 on error



File writing

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t nbyte);
```

■ parameters

- *fd* : file descriptor
- *buf* : buffer address
- *nbyte* : number of bytes to write

■ return

- number of bytes written successfully if OK
- -1 on error



File create/lseek example

create-ex.c

```
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/fcntl.h>
#include <unistd.h>

int main(void)
{
    int fd;
    char buf1[] = "Test1 data";
    char buf2[] = "Test2 data";

    if ((fd == creat ("test.txt", S_IRUSR | S_IWUSR | S_IRGRP |
                     S_IROTH)) < 0) {
        printf("creat error");
        return 1;
    }
    write(fd, buf1, 10);
    if(lseek(fd, 6L, SEEK_SET) == -1) {
        printf("lseek error");
        return 2;
    }
    write(fd, buf2, 10);

    return 0;
}
```

```
$ ls
a.out  test.c
$ ./a.out
$ ls
a.out  test.c  test.txt
$ cat  test.txt
Test1 Test2 data
$
```



File copy example (1)

fcopy2-ex.c

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

#define BUFFER_SIZE 1024

int main(int argc, char *argv[])
{
    int fdi, fdo;
    char buf[BUFFER_SIZE];
    ssize_t n;

    if(argc != 3) {
        perror(argv[0]);
        return 1;
    }
}
```



File copy example (2)

fcopy2-ex.c

```
if((fdi = open(argv[1], O_RDONLY)) == -1) {
    perror(argv[1]);
    return 2;
}

if((fdo = open(argv[2], O_WRONLY | O_CREAT | O_TRUNC |
                O_EXCL, S_IRUSR | S_IWUSR)) == -1) {
    perror(argv[2]);
    return 3;
}
while((n = read(fdi, buf, BUFFER_SIZE)) > 0)
    write(fdo, buf, n);

close(fdi);
close(fdo);
return 0;
}
```

