

System Programming

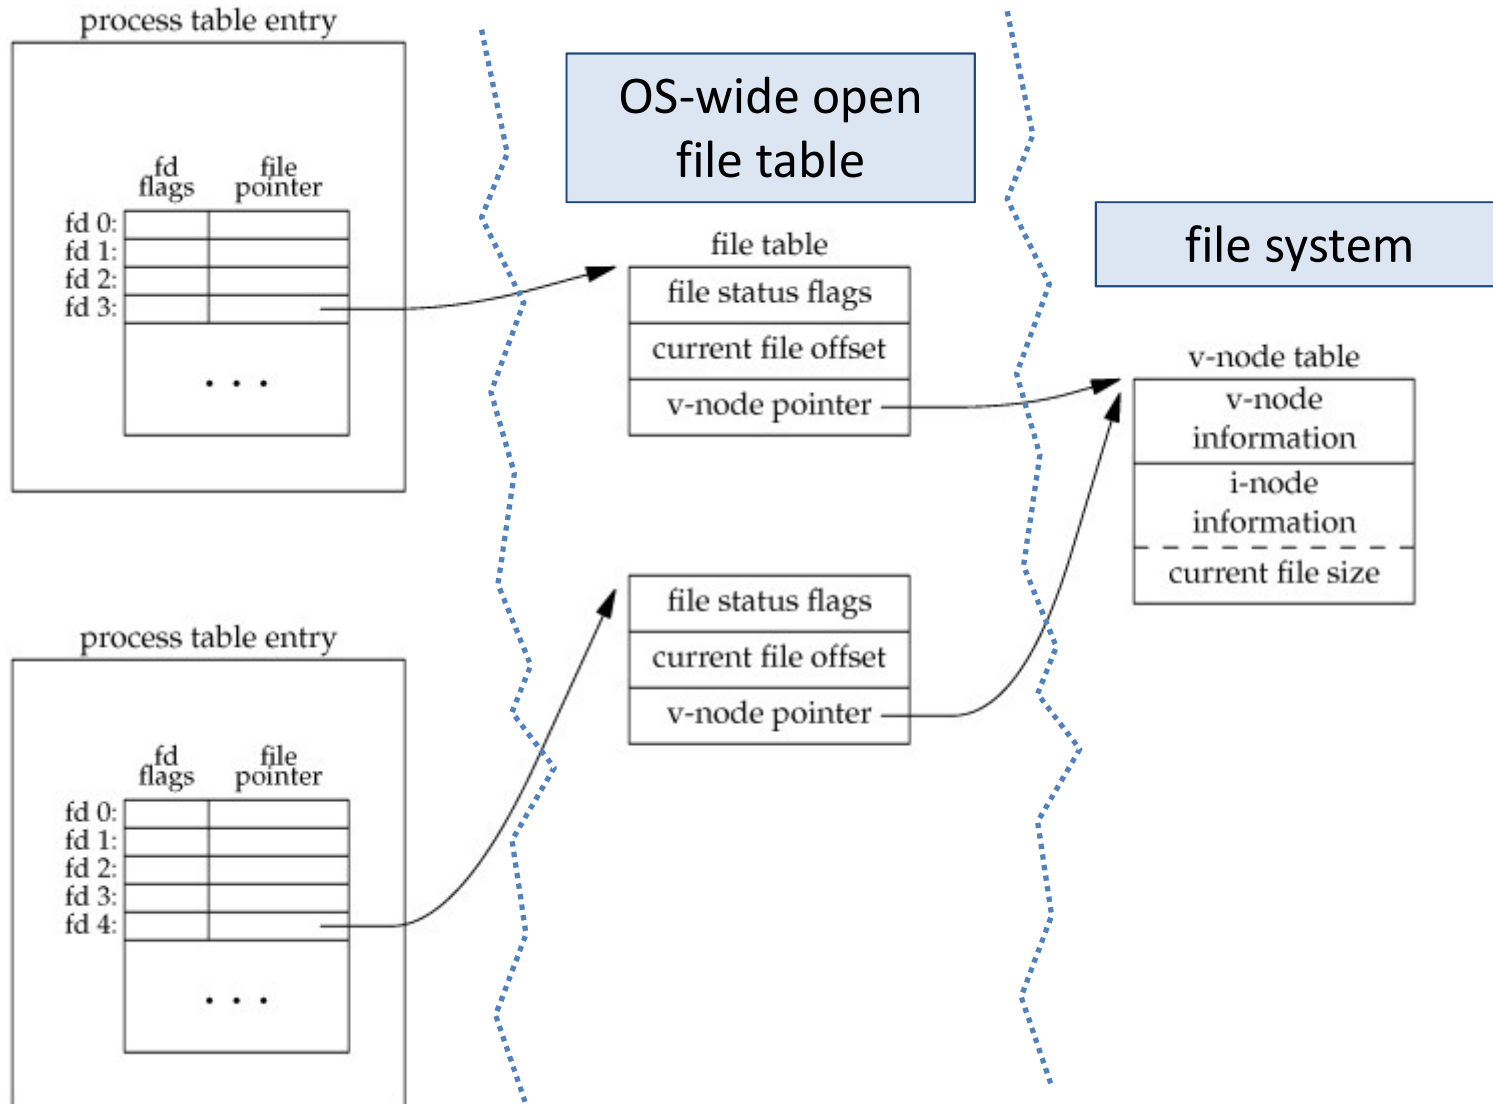
3. File IO (2): System Call

Seung-Ho Lim

Dept. of Computer & Electronic Systems Eng



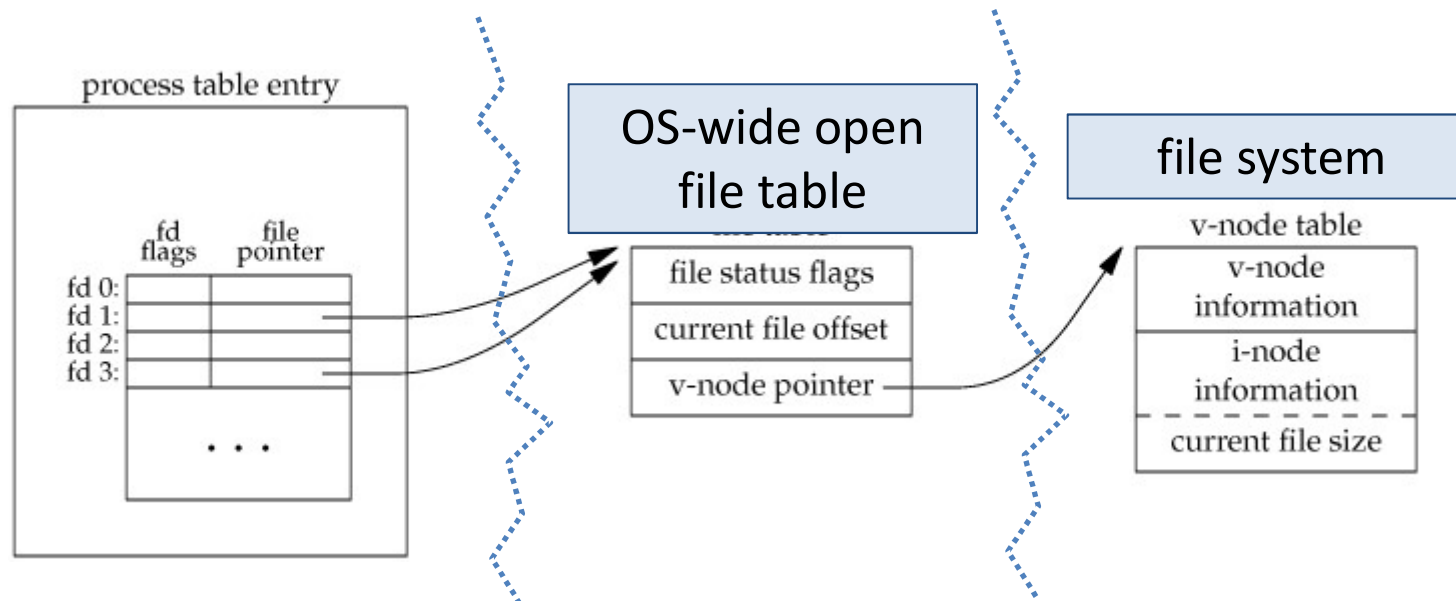
File descriptors & File table



Duplication of file descriptor (1)

```
#include <unistd.h>
int dup(int fd);
```

- parameters
 - *fd* : file descriptor to duplicate
- return
 - newly duplicated file descriptor if OK
 - -1 on error



Duplication of file descriptor (2)

```
#include <unistd.h>
int dup2(int fd1, int fd2);
```

- parameters
 - *fd1* : source file descriptor
 - *fd2* : destination file descriptor
- return
 - copied file descriptor (should be same as *fd2*) if OK
 - -1 on error
- note
 - functionally same as `dup` except that `dup2` designates destination file descriptor (*fd2*) the user wants



I/O redirection example (1)

io-redir.c

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int main()
{
    int backup_des, stdout_des, ofdes;

    stdout_des = fileno(stdout);
    backup_des = dup(stdout_des);
```



I/O redirection example (2)

io-redir.c

```
printf("Hello, world! (1)\n");

ofdes = open("test.txt", O_WRONLY|O_CREAT|O_TRUNC,
             S_IRUSR|S_IWUSR);

dup2(ofdes, stdout_des);
printf("Hello, world! (2)\n");

dup2(backup_des, stdout_des);
printf("Hello, world! (3)\n");

close(ofdes);
}
```



Link (1)

■ Symbolic link

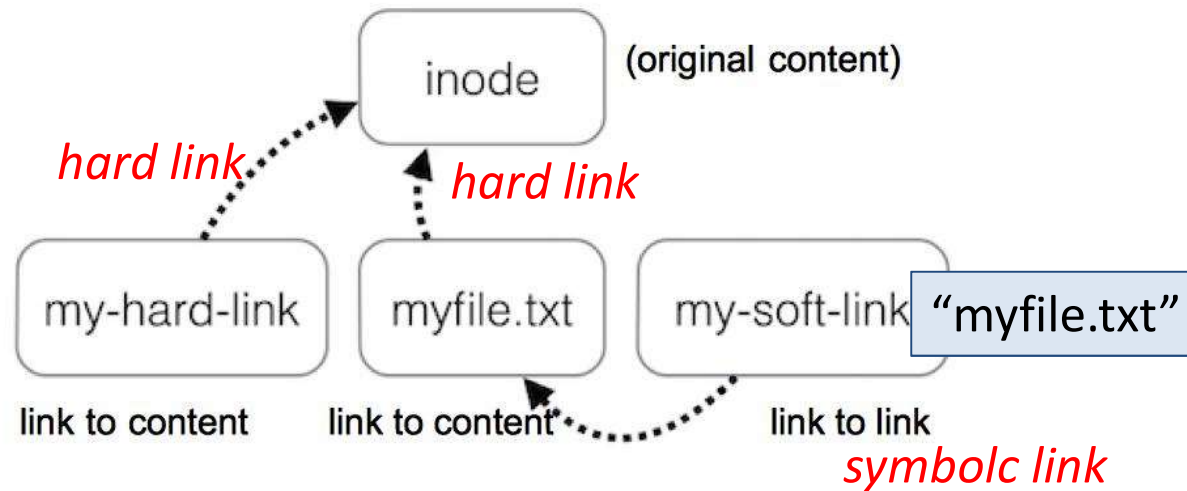
- also called *soft link*
- a **file** which records a path name to a target file
- if the target file is removed, the link is not valid any longer, but the symbolic link file still exists

■ Hard link

- another link which points to an existing inode which is already used by another file
- an inode can be shared by two or more filenames (hardlinks)
- once a file is update, the update can be seen in all the hardlink files
- though a file is removed, another hardlink can access the file



Link (2)



■ Linux command

- symbolic link:
`$ ln -s original_file symbolic_link_name`
- hard link:
`$ ln original_file hard_link_name`



Hard link

```
#include <unistd.h>

int link(const char *existing, const char *new_Link);
```

- parameters
 - *existing* : original file name
 - *new_link* : new link name which will share the inode of existing file
- return
 - 0 if OK
 - -1 on error



Hard link example (1)

hlink-ex.c

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    if(argc != 3) {
        perror("argument error");
        return 1;
    }
    if (link(argv[1], argv[2]) < 0) {
        perror("link fail");
        return 2;
    }
}
```



Hard link example (2)

■ *Run & Results*

```
$ ls -l my*
drwxr--r-x 3  oskernel  oskernel  512  Jul 9 21:58  .
-rw-r--r-  1  oskernel  oskernel   15  Jul 9 21:58  myfile

$ ./a.out  myfile  myhardlink

$ ls -l my*
drwxr-xr-x 3  oskernel  oskernel  512  Jul 9 22:01  .
-rw-r--r- 2  oskernel  oskernel   15  Jul 9 22:01  myfile
-rw-r--r- 2  oskernel  oskernel   15  Jul 9 22:01  myhardlink

$
```



Symbolic link

```
#include <unistd.h>

int symlink(const char *existing, const char *link_name);
```

- parameters
 - *existing* : original file name
 - *link_name* : link name which points the existing file
- return
 - 0 if OK
 - -1 on error



Symbolic link example (1)

symlink-ex.c

```
#include <unistd.h>

int main(int argc, char *argv[])
{
    if(argc != 3) {
        perror("argument error");
        return 1;
    }
    if (symlink(argv[1], argv[2]) < 0) {
        perror("symlink fail");
        return 2;
    }
}
```



Symbolic link example (2)

■ *Run & Results*

```
$ ls -l my*  
-rw-r--r-  1 oskernel oskernel  22  Jul 7 22:21  myfile  
  
$ ./a.out myfile mylink  
  
$ ls -l my*  
-rw-r--r--  1 oskernel oskernel  22  Jul 7 22:21  myfile  
lrwxrwxrwx  1 oskernel oskernel   6  Jul 9 22:18  mylink
```

- *What does the file “mylink” contain?*
 - “myfile” → thus, the file size is 6 bytes.



system calls and symbolic links

system calls which does NOT follow a symbolic link	system calls which follow a symbolic link
lchown, lstat, remove, readlink, rename, unlink these system calls handle a symbolic link itself as a FILE!	accessm chdir, chmod, chown, creat, exec, link, mkdir, mkfifo, mknod, open, opendir, pathconf, stat, truncate

Following a link (1)

```
#include <unistd.h>
```

```
int readlink(const char *path, void *buf, size_t bufsize);
```

- parameters
 - *path* : link name
 - *buf* : buffer address (original file's name)
 - *bufsize* : buffer size
- return
 - number of bytes read if OK
 - -1 on error



Following a link example (1)

readlink-ex.c

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <stdio.h>

#define BUFFER_SIZE 100

int main(int argc, char *argv[])
{
    char buf[BUFFER_SIZE];
    int read_size = 0;

    if(argc != 2) {
        perror("argument error");
        return 1;
    }
    if ((read_size = readlink(argv[1], buf, BUFFER_SIZE)) < 0) {
        perror("readlink");
        return 2;
    }

    buf[read_size] = '\0' ;
    printf("%s\n", buf);
}
```



Following a link example (2)

Run & Results

```
$ ls -ld myfile mylink
-rw-r--r-- 1 oskernel oskernel 22 Jul 7 22:21 myfile
lrwxrwxrwx 1 oskernel oskernel 6 Jul 7 22:18 mylink -> myfile

$ ./a.out mylink
myfile

$
```



File information retrieval (1)

```
#include<sys/types.h>
#include<sys/stat.h>
#include<unistd.h>

int stat(const char *path, struct stat *buf);
```

- parameters
 - *path* : file path name
 - *buf* : address of struct stat (which contains a file information)
- return
 - 0 if OK
 - -1 on error

File information retrieval (2)

```
int lstat(const char *path, struct stat *buf);
```

- basically, same as stat(), but
 - if the file is a symbolic link, retrieve the information of the link file itself (does not follow the link)
- parameters
 - *path* : file path name
 - *buf* : address of struct stat (which contains a file information)
- return
 - 0 if OK
 - -1 on error



File information retrieval (3)

```
int fstat(int fd, struct stat *buf)
```

- parameters
 - *fd* : file descriptor
 - *buf* : address of struct stat (which contains a file information)
- return
 - 0 if OK
 - -1 on error

struct stat *fields*

```
struct stat {
    dev_t    st_dev;        // device
    ino_t    st_ino;        // i-node #
    mode_t   st_mode;       // access mode
    nlink_t  st_nlink;      // number of hard links
    uid_t    st_uid;        // owner id
    gid_t    st_gid;        // group owner
    dev_t    st_rdev;       // device type (if inode device)
    off_t    st_size;       // total size of file
    long     st_blksize;    // block size for I/O
    long     st_blocks;     // number of blocks allocated
    time_t   st_atime;      // time of last access
    time_t   st_mtime;      // time of last modification
    time_t   st_ctime;      // time of last change (including
                           // ownership change)
};
```



Macros for struct stat

Macro	Functions
S_ISREG(st_mode)	return true if the file is <i>regular file</i>
S_ISDIR(st_mode)	return true if the file is <i>directory</i>
S_ISCHR(st_mode)	return true if the file is <i>character device file</i>
S_ISBLK(st_mode)	return true if the file is <i>block device file</i>
S_ISFIFO(st_mode)	return true if the file is <i>FIFO file</i>
S_ISLNK(st_mode)	return true if the file is <i>link file</i>
S_ISCOCK(st_mode)	return true if the file is <i>socket file</i>

File stat example (1)

fstat-ex.c

```
#include <sys/stat.h>
#include <sys/types.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    struct stat statbuf;

    if(argc != 3) {
        perror("argument error");
        return 1;
    }
    if (!strcmp(argv[1], "stat")) {
        if (stat(argv[2], &statbuf) < 0) {
            perror("stat");
            return 2;
        }
    }
}
```



File stat example (2)

fstat-ex.c

```
else if (!strcmp(argv[1], "fstat")) {
    int filedes = open(argv[2], O_RDWR);
    if (fstat(filedes, &statbuf) < 0) {
        perror("stat");
        return 3;
    }
}
else if(!strcmp(argv[1], "lstat")) {
    if (lstat(argv[2], &statbuf) < 0) {
        perror("lstat");
        return 4;
    }
}
if(S_IREG(statbuf.st_mode))
    printf("%s is Regular File\n", argv[2]);
if(S_ISDIR(statbuf.st_mode))
    printf("%s is Directory\n", argv[2]);
if(S_ISLNK(statbuf.st_mode))
    printf("%s is Link File\n", argv[2]);
}
```



File stat example (3)

■ Run & Results

```
$ ls -ld mydir myfile mylink
drwxr-xr-x 2    root root   512      Jul 9 19 : 59    mydir
-rw-r--r-- 1    root root    26      Jul 2 23 : 41    myfile
lrwxrwxrwx 1    root root    12      Jul 9 19 : 59    mylink ->mydir/myfile

$ ./a.out stat myfile
myfile is Regular File

$ ./a.out stat mydir
mydir is Directory

$ ./a.out stat mylink
mylink is Regular File

$ ls -l mydir
-rwxrwxrwx 1 peace peace 0 Jul 9 19 : 59 myfile

$ ./a.out lstat mylink
Mylink is Link File

$ ./a.out fstat mylink
mylink is Regular File
```

