

System Programming

4. Concurrent Processes

Seung-Ho Lim

Dept. of Computer & Electronic Systems Eng



Concurrent Processes (1)

- A running program-system consists of several dynamic concurrent processes
- A set of collaborating processes with IPC
- Concurrent process programming tools
 - Concurrent languages (Ada, Java, PathPascal, Modula II)
 - C, C++,... with system calls for concurrent programming (Linux, UNIX, Win32 Environment)



Concurrent Processes (2)

- Applications (Objectives)
 - Networked (or distributed) applications
 - Network App : client processes, server processes
 - Distributed App : chess program, cloud computing
 - Real-time applications
 - Aircraft control : cockpit display, Flight management,
 - Missile system, Automobile : many sensors, controllers, and servo motors(actuators) are handled by multiple processes
 - Usually, a process for each sensor
 - Parallel applications: 3D graphic, many CPUs, many parallel processes
 - To utilize more CPUs, GPUs
 - CPU/IO overlap in a program level
 - When a process does computing, another process does I/O in a program.
 - Asynchronous event handling
 - Handling of multiple input sources: don't know which device makes input first.
 - In this case, assign a process to each input device.



Process & PID

- Every process has a unique process ID (PID, a non-negative integer)
 - Although they are unique, the IDs are reused after termination
 - PID 0 is usually the scheduler known as the *swapper*
 - PID 1 is usually the **init** process (invoked at the end of the bootstrap procedure, located in `/etc/init` or `/sbin/init`)
 - PID 2 is *pagedaemon* responsible for supporting the paging of the virtual memory system
- Each process has its own PCB(process control block)
 - contains all the information about a process
 - PID, UID, GID, current working directory, terminal, priority, state, CPU usage, memory map, open files and locks, ...



System calls for PIDs

```
#include <unistd.h>

pid_t getpid(void);
// Returns: process ID of calling process

pid_t getppid(void);
// Returns: parent process ID of calling process

uid_t getuid(void);
// Returns: real user ID of calling process

uid_t geteuid(void);
// Returns: effective user ID of calling process

gid_t getgid(void);
// Returns: real group ID of calling process

gid_t getegid(void);
// Returns: effective group ID of calling process
```



Process creation in Linux

```
#include <unistd.h>

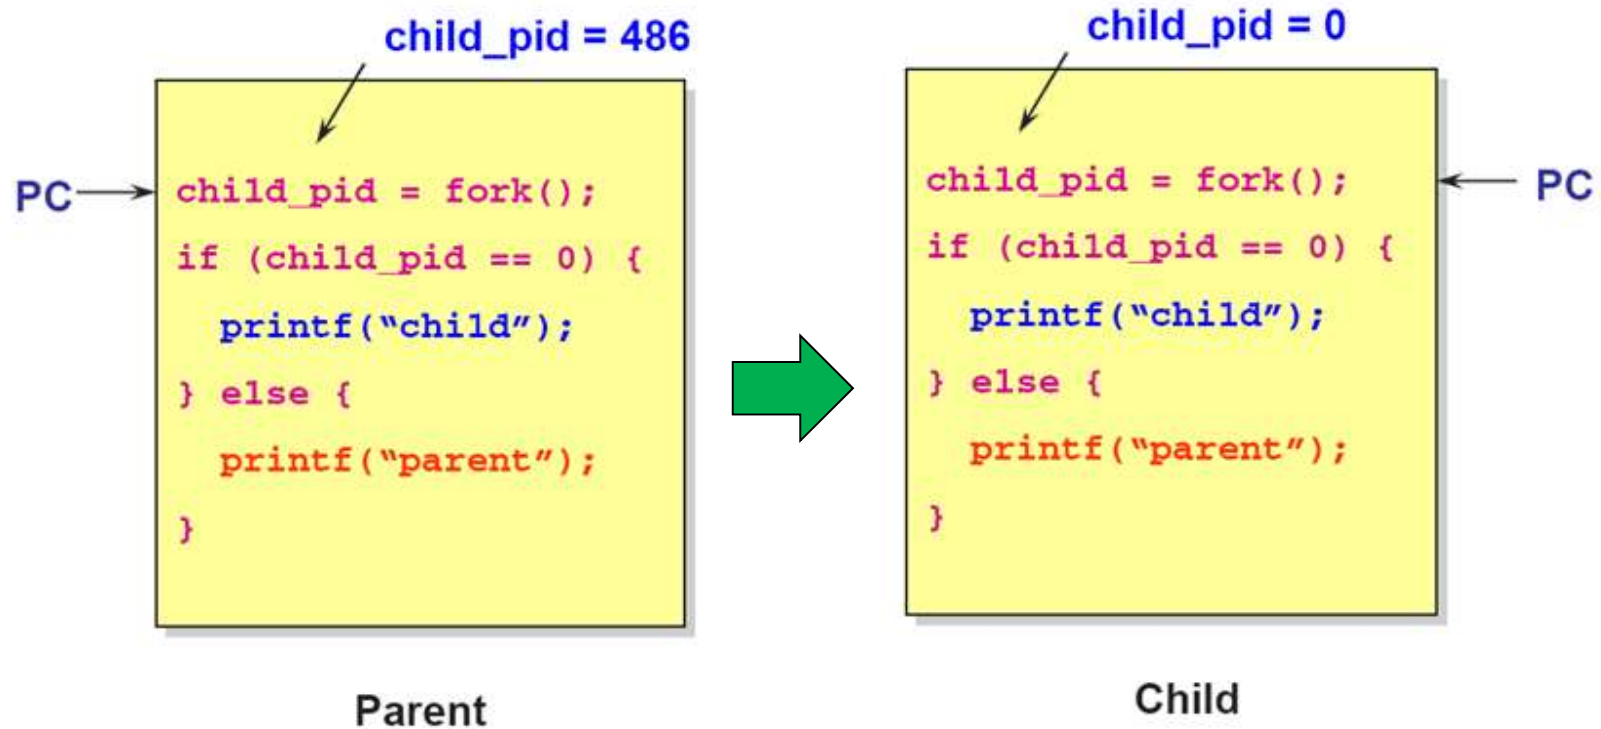
pid_t fork(void);
```

- creates a new process called *child process*
 - the child is a clone (copy) same as the parent
 - the child starts by returning from this fork() call (Why?)
- returns
 - PID of the new born child process in parent process
 - 0 in the child process
 - -1 on error



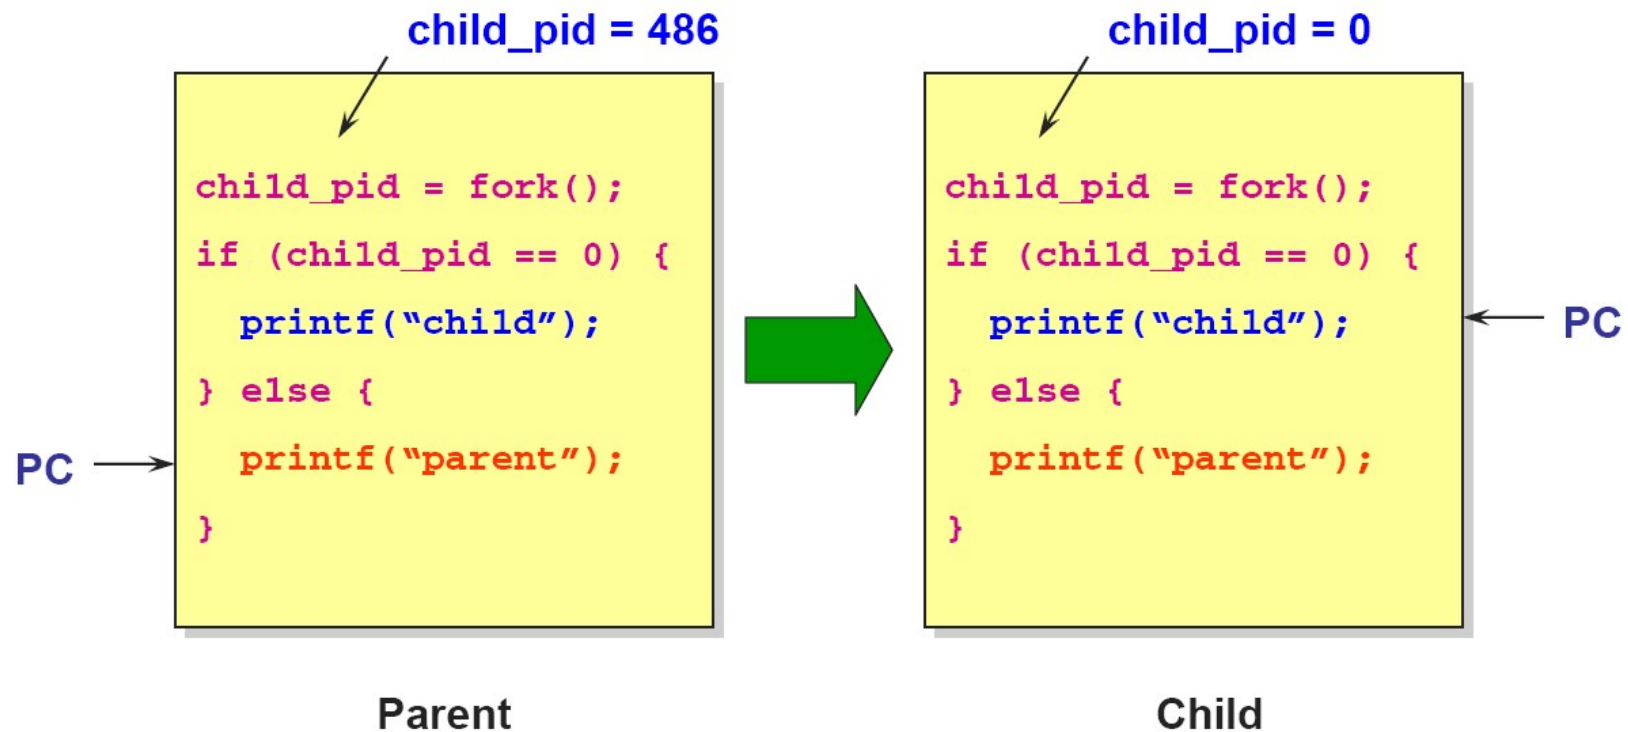
fork() example (1)

- Duplicating the address spaces




fork() example (2)

- Divergence of their execution flow



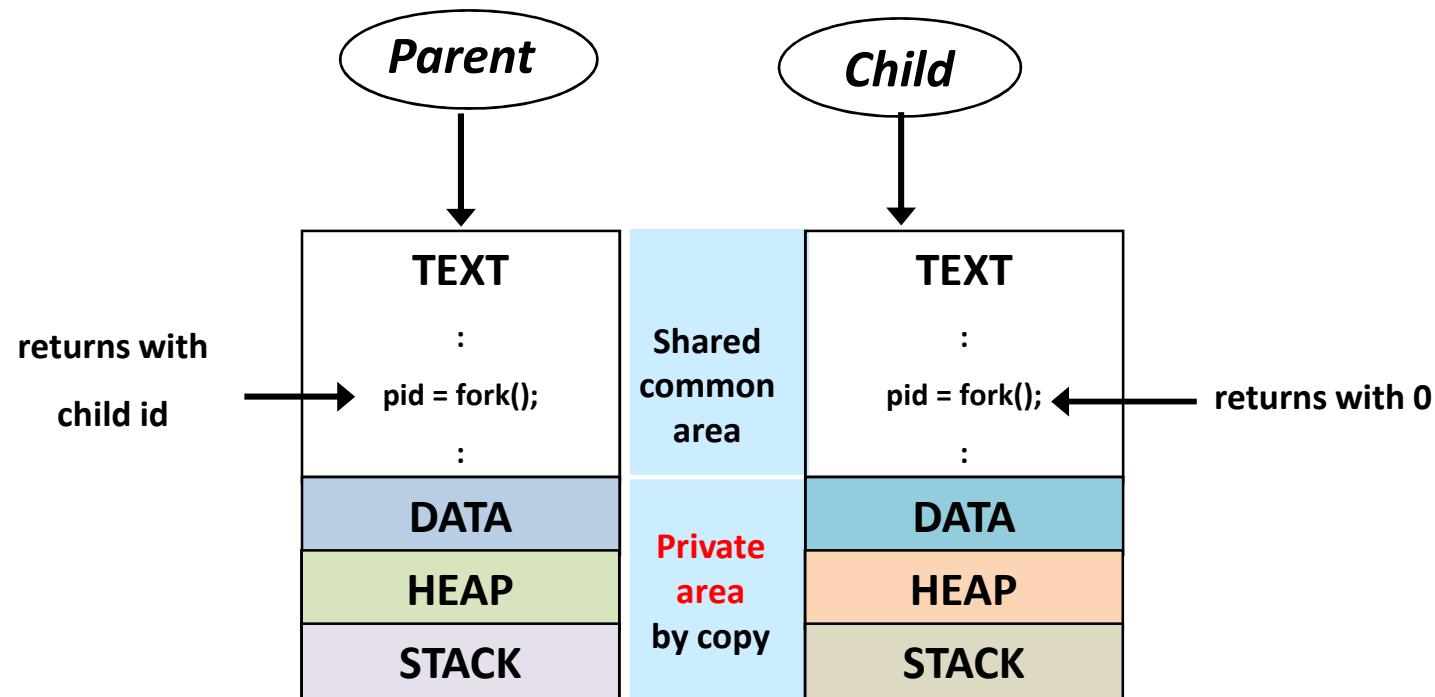
Parent and Child (1)

- Child is a copy (or clone) of the parent process
- Child inherits most resources of the parent at the fork() call
 - share the same text (program code and constant data)
 - real/effective user-id, group-id,
 - current working, root directory,
 - **open files** before fork including tty (stdin, stdout, stderr),
 - share **r/w offsets** of the files which the parent opened before fork.



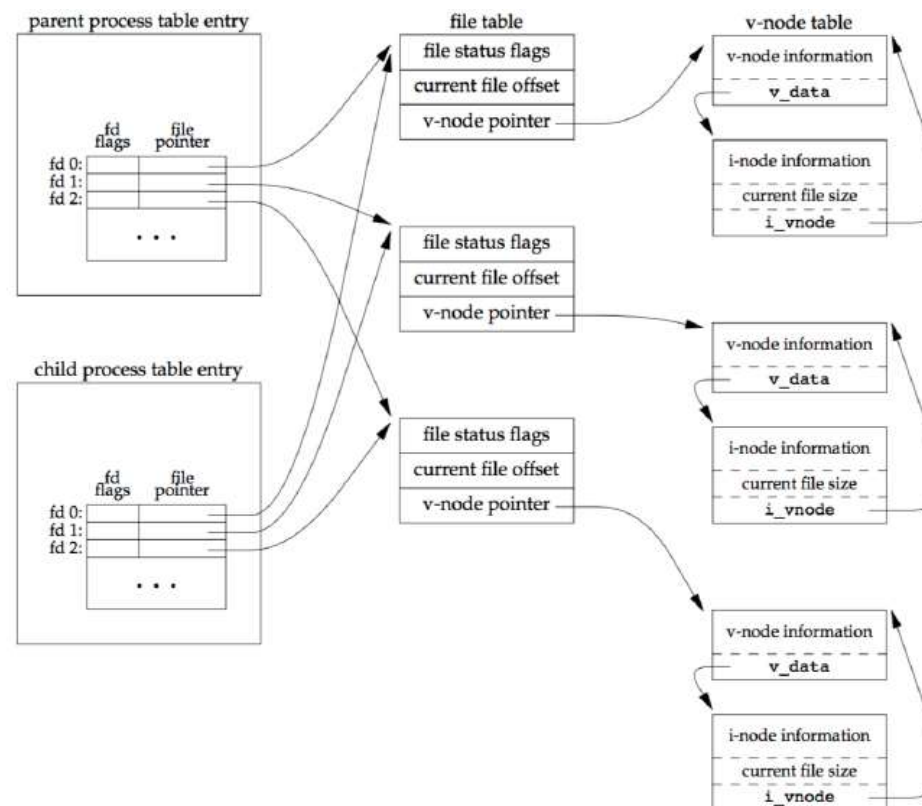
Parent and Child (2)

- But, in reality, those are different processes
 - different PID, different PCB,
 - scheduled independently (concurrently) by the kernel
 - each has some private resources
 - **different data/heap/stack** (copied when forked)



Parent and Child (3)

- the fork() may fail
 - if too many processes are already in the system
 - if the total number of processes for this real user ID exceeds the system's limit (CHILD_MAX)



Terminating a process (1)

```
#include <stdlib.h>
```

```
void exit(int status);
```

- `exit()` function sends **status** to its parent
 - the status informs how the child process has been terminated
 - its parent gets the status value via `wait()` call (see later!)
- A process can terminate normally in five ways
 - executing a return from the main function.
 - calling the `exit()` function.
 - calling the `pthread_exit` function from the last thread in the process.



Terminating a process (1)

- Kernel code executes these for a terminated process
 - the code closes all the open descriptors for the process
 - releases the memory allocated

cf. the terminated process remains in the system until its parent returns from `wait()` system call

- called zombie process in Linux

Waiting a child process (1)

```
#include <sys/wait.h>
pid_t wait(int *status);
```

- caller (usually parent) will be blocked until any child terminates
 - when resumes, removes the PCB of the child.(a funeral)
- return
 - terminated PID if Ok
 - -1 on error
- parameters
 - int *status : variable address where the status will be stored



Waiting a child process (2)

```
#include <sys/wait.h>
pid_t waitpid(pid_t pid, int *status, int options);
```

- caller (usually parent) will be blocked until the designated process (pid) terminates
 - same as wait() except the pid designation
 - how to make wait() do as waitpid()?
- parameters
 - pid
 - if pid > 0 waits for the child whose process ID equals pid
 - if pid == -1 waits for any child process, waitpid is equivalent to wait
 - options
 - WNOHANG: for non-blocking call



Macros for exit status

- **WEXITSTATUS(status)**
 - fetch the exit code which the child sends
- **WIFEXITED(status)**
 - true if the child terminated normally
- **WIFSIGNALED(status)**
 - true if status was child terminated abnormally by a signal
- **WIFSTOPPED(status)**
 - true if status was returned for a child that is currently stopped
- **WTERMSIG(status)**
 - fetch the signal number that caused the child to be terminated
- **WCOREDUMP(status)**
 - true if a core file or the terminated process was generated



wait() and exit() example

```
#include <wait.h>

main()
{
    int res, pid ;

    res = fork() ;

    if (res == 0) {
        // do something
        exit(77); // 77 is an example exit status
    }
    else {
        pid = wait(&status); // waiting until the child terminates
        printf("the child %d is exited with %d \n",
            pid, WEXITSTATUS(status));
    }
}
```

parent process

```
#include <wait.h>

main()
{
    int res, pid ;

    res = fork() ;

    if (res == 0) {
        // do something
        exit(77); // 77 is an example exit status
    }
    else {
        pid = wait(&status); // waiting until the child terminates
        printf("the child %d is exited with %d \n", pid,
            pid, WEXITSTATUS(status));
    }
}
```

child process



exec() system call

- When a process calls one of the exec functions, that process is completely **replaced** by the new program
 - the new program **starts** executing **at its main function**
 - PID does not change across an exec
 - exec merely replaces the current process (text, data, heap, and stack segment)
- The program body is changed but the followings remain
 - priority
 - process id (same process), working directory
 - time left until an alarm
 - pending signals
 - open files and file locks



exec family (1)

```
#include <unistd.h>
```

```
int execl(const char *pathname, const char *arg0, ... /* (char *)0 */ );
```

```
int execv(const char *pathname, char *const argv[]);
```

```
int execle(const char *pathname, const char *arg0, ... /* (char *)0, char *const  
envp[] */ );
```

```
int execve(const char *pathname, char *const argv[], char *const envp[]);
```

```
int execlp(const char *filename, const char *arg0, ... /* (char *)0 */ );
```

```
int execvp(const char *filename, char *const argv[]);
```

```
int fexecve(int fd, char *const argv[], char *const envp[]);
```

```
// All return: -1 on error, no return on success
```



exec family (2)

■ Difference

Function	pathname	filename	fd	arg list	argv[]	environ	envp[]
execl	•			•		•	
execlp		•			•	•	
execle	•			•			•
execv	•				•	•	
execvp		•			•	•	
execve	•				•		•
fexecve			•		•		•
letter		p	f	l	v		e

■ letters


- *v* in its name, argv's are a vector: `const * char argv[]`
- *l* in its name, argv's are a list: `const * char arg0, ...`
- *e* in its name, takes environment variables: `char * const envp[]`
- *p* in its name, PATH environment variable to search for the file



exec() call example


```
main()
{
    pid = fork() ;
    if (pid==0)
        exec("test.exe");
}
```

parent process

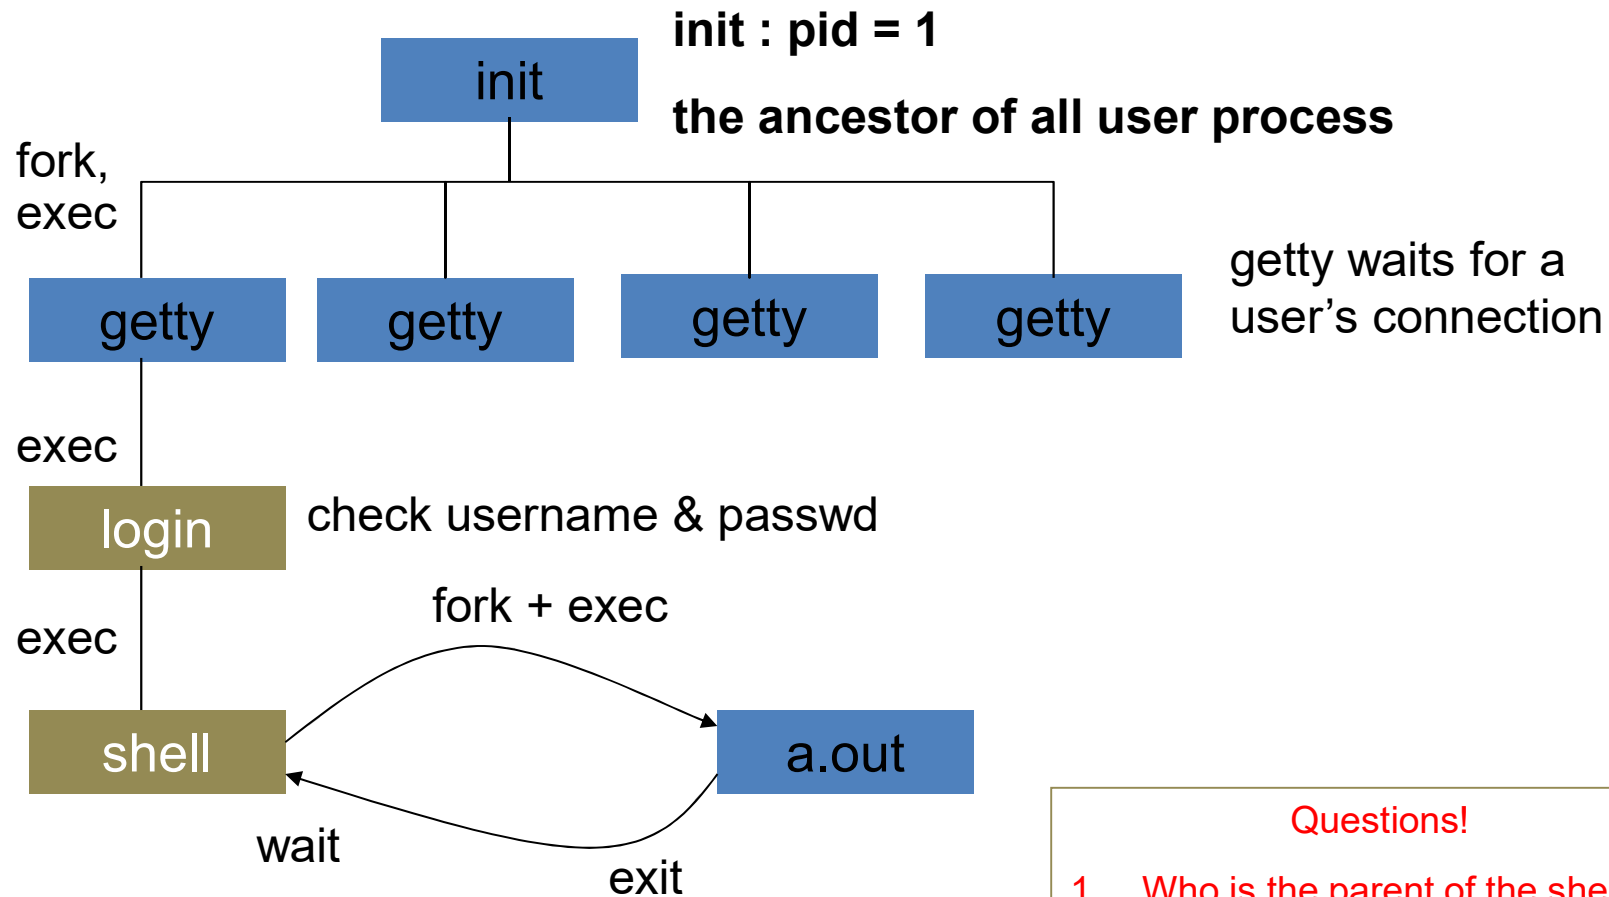


```
main()
{
    // test.exe program
    main()
    {
        // this program replace all address
        // space of the forked process.
        printf("Hello, world! \n");
    }
}
```

child process



User process tree



Questions!

1. Who is the parent of the shell?
2. Who is the parent of the a.out?

Shell & process tree

\$: shell, the process of this user

\$./a.out : shell and a.out : processes of this user

\$ cat myfile : shell and cat : processes of this user

\$./a.out & : a.out as a background process

\$ rm myfile : shell, a.out and rm are running processes

*\$./a.out > output : stdout is changed to “output” before exec().
This is called “I/O redirection”.*



Simple *shell* example

```
while (not logout) {
    print prompt;
    get a line string;
    parse the line; // a.out, a.out&, ...
    if ((pid = fork()) == 0) { // child
        ...           // IO redirection if designated
        execv("parse name", argv);
    }
    if (foreground) {
        while (pid != wait()); // wait for the death of the foreground process
        // if a background did exit(), it is removed here as a side effect.
    } // if background, does not wait()
}
// You can make your own shell !
```

