

System Programming

7. Record Lock

Seung-Ho Lim



Dept. of Computer & Electronic Systems Eng.



Concurrent-Readers / Exclusive-Writers

■ Problem Definition

- Some **writer & reader processes** access a shared resource (variables, files) concurrently.
- **Mutual exclusion** must be provided to readers & writers
 - one by one accesses → low performance
 - Readers can share a shared resource because they do not modify it !
- An improvement
 - Any writer process must access the shared resource exclusively from other writers and readers. (**writing -> exclusive lock**)
 - But when reading, a reader must be protected from other writer processes, however, other readers can join in concurrent reading. (to enhance the performance! 80% of DB operations are readings!) (**reading -> shared lock**)



Concurrent-Readers/ Exclusive-Writers Problem

■ Writer's lock (Exclusive lock)

- **Mutex** for writers
- Same as the original **mutex**
- When writing, the writing is protected from all other writers & readers

■ Reader's lock (Shared lock)

- A special **mutex** for readers.
- At the time when a reader requests the reader's lock, the reader can join if previous readers are currently accessing the shared resource. (reader's lock is already set) -> shares the reader's lock
- At the time when a reader requests the reader's lock, the reader must wait if a previous writer is accessing the shared resource. (writer's lock is already set)

Record Lock in a File (1)

```
#include <fcntl.h>
```

```
int fcntl(int fildes, int cmd, struct flock *lock);
```

- parameters

- fildes : file descriptor
- cmd : command
 - F_GETLK : check if the lock can be acquired
 - » if already locked by someone, return a filled lock structure
 - » if it can be acquired, return a lock structure with F_UNLCK (l_type)
 - F_SETLK : try to set lock as designated in lock argument.
 - » called with l_type == F_RDLCK or F_WRLCK
 - » if already locked, return -1 immediately
 - » on an error, errno = EACCESS or EAGAIN



Record Lock in a File (2)

- **F_SETLKW** :

- » blocking version for F_SETLK

- » if already locked, the calling process must be blocked until it can get the lock.

- F_SETLK and l_type == F_UNLCK

- » release a lock designated by the flock structure



Record Lock in a File (3)

```
struct flock {  
    .....  
    short      l_type; // type of lock: F_RDLCK, F_WRLCK, F_UNLCK  
                  // F_RDLCK : reader's lock  
                  // F_WRLCK : writer's lock  
                  // F_UNLCK : unlock  
    short      l_whence; // SEEK_SET, SEEK_CUR, SEEK_END  
  
    off_t       l_start; // lock start position  
    off_t       l_len;   // lock region length (0 for entire file)  
    pid_t       l_pid;   // pid of a process that has a lock  
                  // (used only in F_GETLK)  
}
```



Example: Record Processing (1)

rec-proc.c

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>

# define NUM_RECORDS 100

struct record {
    char name[20];        // account owner
    int  id;              // account number
    int  balance;
};

void get_new_record(struct record *curr);
void display_record(struct record *curr);
```



Example: Record Processing (2)

```
int main( int argc, char *argv[])
{  FILE *fp; struct record current;
  int record_no;
  int fd;
  long pos;
  char yes;

  fd = open( argv[1], O_RDWR | O_CREAT , S_IRUSR | S_IWUSR);

  if(( fp = fdopen( fd, "r+")) == NULL)
  {  perror( argv[1]);
    exit(2);
  }
  printf("enter record number: ");
  scanf("%d",&record_no);
  getchar();
```



Example: Record Processing (3)

```
while (record_no >= 0 && record_no < NUM_RECORDS) {
    pos = record_no * sizeof(struct record);
    fseek( fp, pos, SEEK_SET);
    fread(&current, sizeof(struct record), 1, fp);
    display_record (&current);
    printf("update records? yes = y\n");
    scanf("%c", &yes); getchar();
    if (yes == 'y') {
        get_new_record(&current);
        fseek( fp, pos, SEEK_SET);
        fwrite( &current, sizeof( struct record), 1, fp);
        printf("update done\n");
    }
    printf("enter record number: ");
    scanf("%d",&record_no);
    getchar();
}
fclose(fp);
}
```



Example: Record Processing (4)

```
void get_new_record(struct record *curr)
{
    printf("> id? ");
    scanf("%d", &curr->id);
    printf("> name? ");
    scanf("%s", curr->name);
    printf("> balance? ");
    scanf("%d", &curr->balance);
}

void display_record(struct record *curr)
{
    printf("\n");
    printf("id: %d \n", curr->id);
    printf("name: %s \n", curr->name);
    printf("balance: %d \n", curr->balance);
    printf("\n");
}
```

Is this code
working good?



Critical Section

Process A

```
read balance;  
update balance(deposit 5);  
write balance;
```

Process B

```
read balance;  
update balance(deposit 6);  
write balance;
```

Think the following scenario!

read by Process A; (current balance: 10)

read by Process B; (current balance: 10)

write by Process A; (update balance: 15)

write by Process B; (update balance: 16)



Record Locking

- Deposit / withdraw (*update balance*)
 - This (write) operation must be done exclusively with other deposit and inquiry operations. This is called a **writer's lock** or an **exclusive lock**.
- Inquiry (*read balance*)
 - This operation can be done with other inquiry operations but must be done exclusively with other write operations.
 - This is called a **reader's lock** or a **shared lock**.

		Request for read lock	Request for write lock
Region currently has	No locks	OK	OK
	One or more read locks	OK	denied
	One write lock	denied	denied



Pseudo Code

```
open the master account file;
while (true) {
    get the operation type and information on the account
    switch (operation) {
        case create:
            get a writer's lock on the record;
            get id and name of the user;
            reset the account information;
            release the writer's lock;

        case inquiry:
            get a reader's lock on that record;
            (if a write lock exists, blocked)
            display the account information;
            release the reader's lock;

        case deposit (or withdraw):
            get a writer's lock on that record;
            (if a writer's or a reader's lock exists, blocked)
            display the account information;
            release the writer's lock;

    }
}
```



Example (1)

reclock.c

```
#include <stdio.h>
#include <stdlib.h>
#include <fcntl.h>
#include <unistd.h>

#define TRUE    1
#define FALSE   0

#define NUM_RECORDS 100

struct record{
    char name[20];    // account owner
    int id;           // account umber
    int balance;      // balance
};

int reclock (int fd, int recno, int len, int type);
void display_record(struct record *curr);
```



Example (2)

```
int main()
{
    struct record current;
    int record_no;
    int fd, pos, i, n;
    char yes;
    char operation;
    int amount;
    char buffer[100];
    int quit=FALSE;

    fd = open("./account", O_RDWR | O_CREAT, S_IRUSR | S_IWUSR);
    while (1) {
        printf("enter account number (0-99): ");
        scanf("%d", &record_no);
        fgets(buffer, 100, stdin);

        if (record_no < 0 && record_no >= NUM_RECORDS)
            break;
        printf("enter operation name (c/r/d/q): ");
        scanf("%c", &operation);
```



Example (3)

```
switch (operation) {
    case 'c' : // create
        reclock(fd, record_no, sizeof(struct record), F_WRLCK);
        pos = record_no * sizeof(struct record);
        lseek(fd, pos, SEEK_SET);
        printf("> id ? ");
        scanf("%d", &current.id);
        printf("> name ? ");
        scanf("%s", current.name);
        current.balance = 0 ;
        n = write(fd, &current, sizeof(struct record));
        display_record (&current);
        reclock(fd, record_no, sizeof(struct record), F_UNLCK);
        break;
    case 'r' : // inquiry
        reclock(fd, record_no, sizeof(struct record), F_RDLCK);
        pos = record_no * sizeof(struct record);
        lseek(fd, pos, SEEK_SET);
        n = read(fd, &current, sizeof(struct record));
        display_record (&current);
        reclock(fd, record_no, sizeof(struct record), F_UNLCK);
        break;
```



Example (4)

```
case 'd' :           // deposit
    reclock(fd, record_no, sizeof(struct record), F_WRLCK);
    pos = record_no * sizeof(struct record);
    lseek(fd, pos, SEEK_SET);
    n = read(fd, &current, sizeof(struct record));
    display_record (&current);
    printf("enter amount\n");
    scanf("%d", &amount);
    current.balance += amount;
    lseek(fd, pos, SEEK_SET);
    write(fd, &current, sizeof(struct record));
    reclock(fd, record_no, sizeof(struct record), F_UNLCK);
    break;

case 'q' :
    quit = TRUE;
    break;

default :
    printf("illegal input\n");
    continue;

}

}

close(fd);
fflush(NULL);

}
```



Example (5)

```
int reclock (int fd, int recno, int len, int type)
{
    struct flock fl;
    switch (type) {
        case F_RDLCK:
        case F_WRLCK:
        case F_UNLCK:
            fl.l_type = type;
            fl.l_whence = SEEK_SET;
            fl.l_start = recno * len;
            fl.l_len = len;
            fcntl (fd, F_SETLKW, &fl);
            return 1;
        default:
            return -1;
    }
}

void display_record(struct record *curr)
{
    printf("\n");
    printf("id: %d \n", curr->id);
    printf("name: %s \n", curr->name);
    printf("balance: %d \n", curr->balance);
    printf("\n");
}
```



File Lock

```
#include <fcntl.h>

int flock(int fd, int operation);
```

- locks entire file (not a record or a part)
- parameters
 - *fd* : open file descriptor
 - operation
 - LOCK_SH: place a shared lock
 - LOCK_EX: place an exclusive lock
 - LOCK_UN: remove an existing lock held by this process
 - can be OR'd with LOCK_NB (non-blocking)



Lock inheritance & release

■ Inheritance

- a lock is not inherited to a child process by `fork()`
- a lock is inherited to a process by `exec()`

■ Release

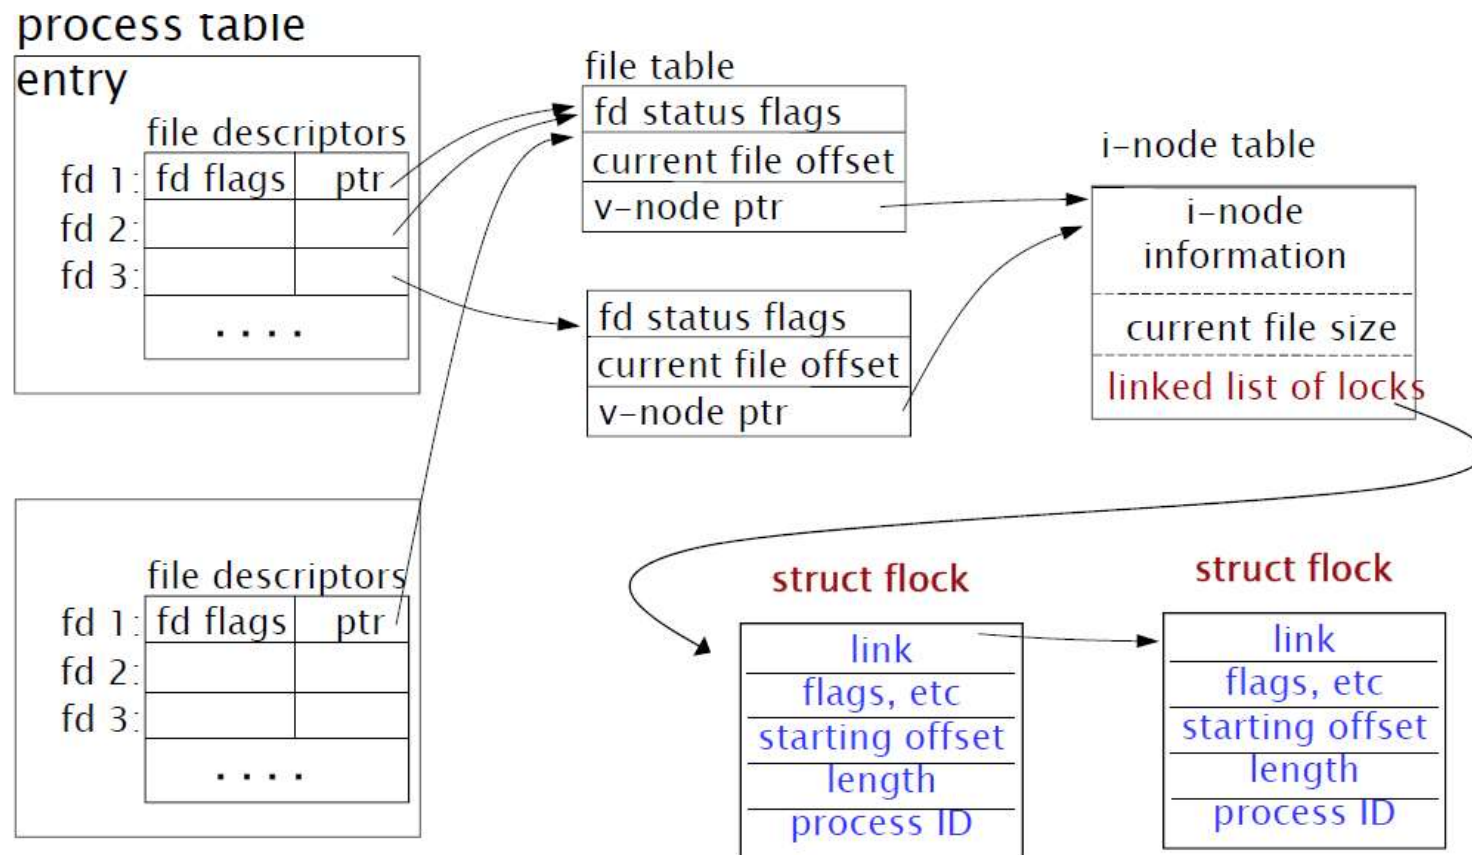
- When a process terminates, all locks created by the process will be released
- When a file descriptor to a file is close, all locks of the file is released

```
fd1 = open("a", ... );  
read_lock(fd1, ... );  
fd2 = dup(fd1);  
close(fd2)
```

```
fd1 = open("a", ... );  
read_lock(fd1, ... );  
fd2 = open("a", ... );  
close(fd2);
```



Flocks in a kernel



- lock is associated with a file (not a fd)
- if any one of fd1, fd2, and fd3 is closed, all the locks will be released



Advisory vs. Mandatory Locking

■ Advisory locking

- locking is not enforced by a kernel
- reads or writes can violate the locking protocol
- processes should voluntarily conform the locking protocol

■ Mandatory locking

- kernel enforces the locking protocol
- any read or write can not violate the locking protocol
- kernel overhead is quite high since it should inspect every read and write call
- how to set the mandatory lock bit

– turn ON the set-group-ID bit and OFF group-execute bit

e.g.

```
$ chmod 2644 lockfile (or chmod g+s,g-x lockfile)
```

```
-rw-r-1r-- 1 rheys 5 Jul 18 12:11 lockfile
```



Ex1

- Producer & Consumer in Pthreads :
 - In the lecture note, an array has been used as the buffer between the producer & consumer threads.
 - We provide linked list, the “*struct LinkedList*” has
 - *Node Head & tail pointers, number_of_items,*
 - *Node : integer data, Node pointer to next*
 - *intertItem: at head, if $number_of_items \geq 100$, then wait,*
 - *getItem: at tail, if $number_of_items == 0$, then wait,*
 - *isEmpty: check $number_of_items == 0$*
 - *isFull: check $number_of_items == 100$*

Ex1

- Producer & Consumer in Pthreads :
 - Modify the program as follows:
 - Buffer array -> integer buffer queue in a linked list.
 - Share linked list between producer & consumer
 - Synchronization with mutex & condition variable
 - Producer
 - » If the linked list is full then wait,
 - » Insert random integer data at the tail of the list if possible.
 - » usleep some amount of time
 - Consumer
 - » If the linked list is empty then wait,
 - » Remove and get the data at the head of the list if possible.
 - » print the data to the console.
 - » usleep some amount of time



EX2

■ File RW lock handling

- Add “withdraw” & “transfer” functions to the example ATM program.
- Commands
 - Existing command : (c/r/d/q)
 - Add command w,t
 - “Withdraw”, “Deposit”
 - » Use a writer’s lock on the record.
 - “Transfer”
 - » Withdraw + deposit must be done atomically.
 - » So
 - » Writer’s lock on the withdraw account;
 - » Writer’s lock on the deposit account;
 - » Do withdraw & deposit;
 - » Unlock all.



Due Date

- Until 5/7
- Submit to eclass