

System Programming

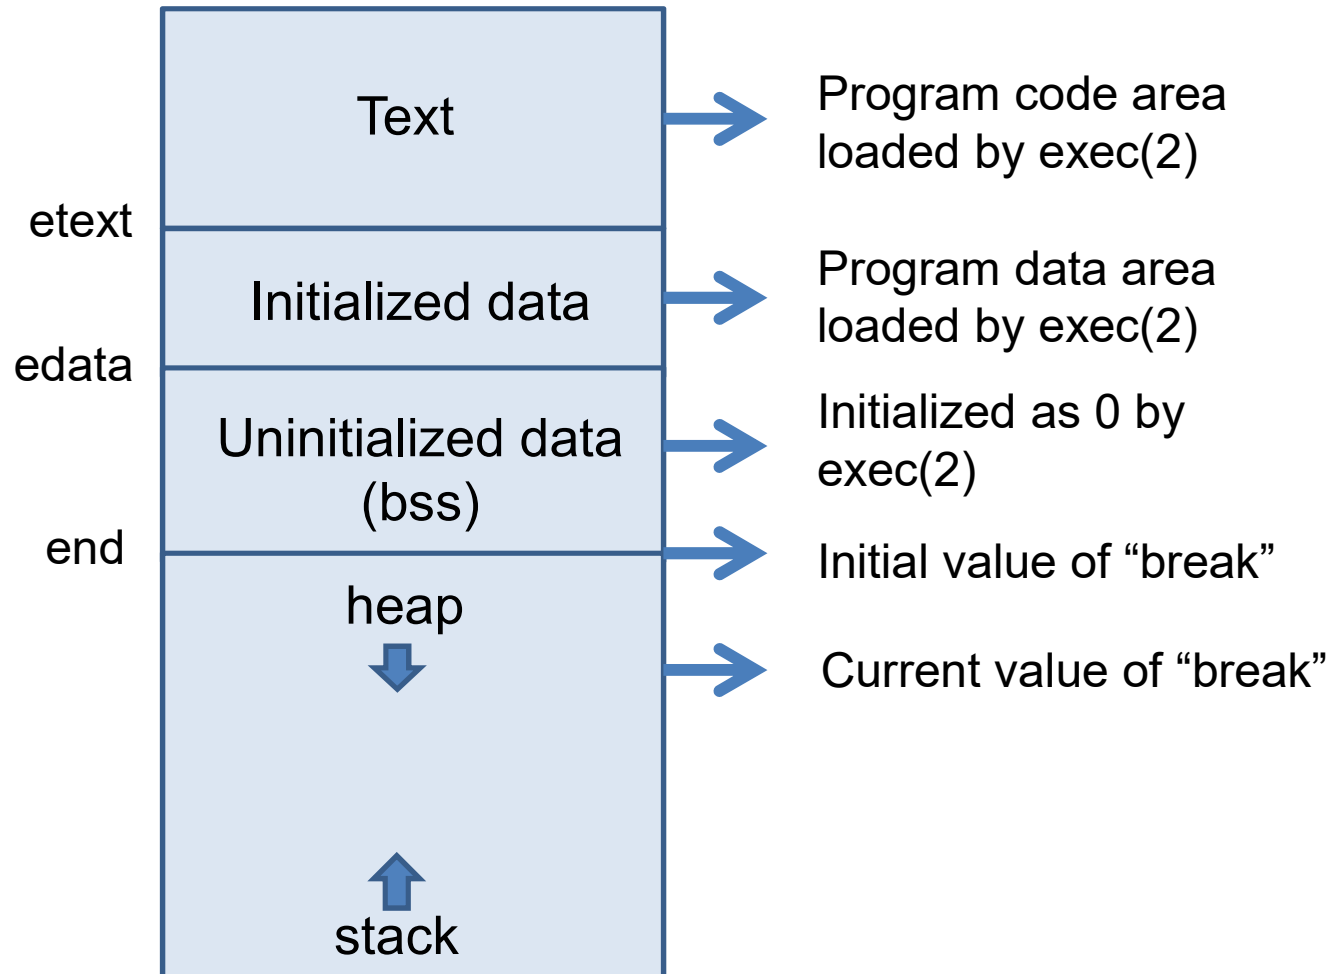
9. Memory

Seung-Ho Lim

Dept. of Computer & Electronic Systems Eng.



Process Address Space



Process Address Space

- **Text Area**
 - CPU instructions of a program
 - Read-only constants
- **Data Area**
 - Global initialized data area + global uninitialized data area (***bss = block started by symbol***)
- **Heap Area**
 - For dynamic memory allocation (malloc(), new, etc.)
- **Stack Area**
 - automatic variables in a function
 - Call frames (arguments, return address, etc.)



Memory for main()


■ vars.c

```
#include <stdio.h>
#include <stdlib.h>
int a = 1;
static int b = 2;

int main() {
    int c = 3;
    static int d = 4;
    char *p;

    p = (char *) malloc(40);
    fun(5);
}

void fun(int n)
{
    int m = 6;
    ...
}
```



Memory for fun()

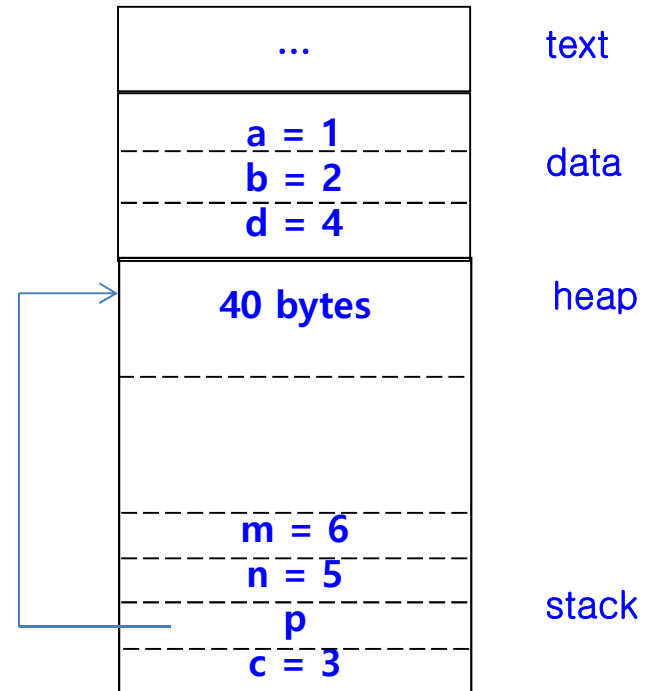
■ vars.c

```
#include <stdio.h>
#include <stdlib.h>
int a = 1;
static int b = 2;

int main() {
    int c = 3;
    static int d = 4;
    char *p;

    p = (char *) malloc(40);
    fun(5);
}

void fun(int n)
{
    int m = 6;
    ...
}
```



Dynamic Memory Allocation

- Reasons to use dynamic allocation
 - This saves memory by requesting and using as much memory as needed when needed.
- `malloc()`
- `calloc()`
- `realloc()`
- `free()`



Memory Allocation

```
#include <stdlib.h>
void *malloc(size_t size);

void free(void *ptr);
```

- Allocate memory space of size bytes and return the starting address in void * type.
- parameters:
 - **size** : Allocate memory space of size bytes
- return:
 - return the starting address in void * type




Memory Allocation usage

```
char *ptr;  
ptr = (char *) malloc(40);
```




```
int *ptr;  
ptr = (int *) malloc(10 * sizeof(int));
```



Memory Allocation usage

```
struct student {  
    int id;  
    char name[10];  
};  
struct student *ptr;  
ptr = (struct student *) malloc(sizeof(struct student));
```



```
struct student *ptr;  
ptr = (struct student *) malloc(n * sizeof(struct student));
```



Memory allocation ex1

■ stud1.c

```
#include <stdio.h>
#include <stdlib.h>
struct student {
    int id;
    char name[20];
};

int main()
{
    struct student *ptr;
    int n, i;
    printf("How many student ? ");
    scanf("%d", &n);
    if (n <= 0) {
        fprintf(stderr, "errpr: wrong number.\n");
        fprintf(stderr, "terminate program\n");
        exit(1);
    }
}
```



Memory allocation ex1

■ stud1.c

```
ptr = (struct student *) malloc(n * sizeof(struct student));
if (ptr == NULL) {
    perror("malloc");
    exit(2);
}

printf("enter student number and name for %d. students\n", n);
for (i = 0; i < n; i++)
    scanf("%d %s\n", &ptr[i].id, ptr[i].name);

printf("\n* student information *\n");
for (i = n-1; i >= 0; i--)
    printf("%d %s\n", ptr[i].id, ptr[i].name);

printf("\n");
exit(0);
}
```



Memory Allocation

```
#include <stdlib.h>
void *calloc(size_t n, size_t size);

#include <stdlib.h>
void *realloc(void *ptr, size_t newsize);
```

- **calloc** : Allocate n memory spaces of size size. Initialize all values to zero. On failure, NULL is returned.
- **Parameter**
 - n : number of spaces to be allocated
 - size : size of each space
- **realloc** : Change the size of the memory allocated by ptr to newsize.
- **Parameters**
 - ptr : pointer for the allocated memory
 - newsize : size for new allocation



Memory Mapped File

■ Objectives

- By mapping a memory area (e.g. *struct*) into a file,
 - **Auto-saving** of memory variables into a file at run time.
 - Memory-reading = reading from a file
 - Memory-writing = writing to a file (to **page cache**)
 - Handle a file data as a memory array or using a pointer var.
- No copy b/w library buffer and kernel buffer
 - improves the file I/O performance
- When several processes map their memory areas into the same file,
 - **Inter-process communication is possible** by reading and writing memory area



Memory Mapped File

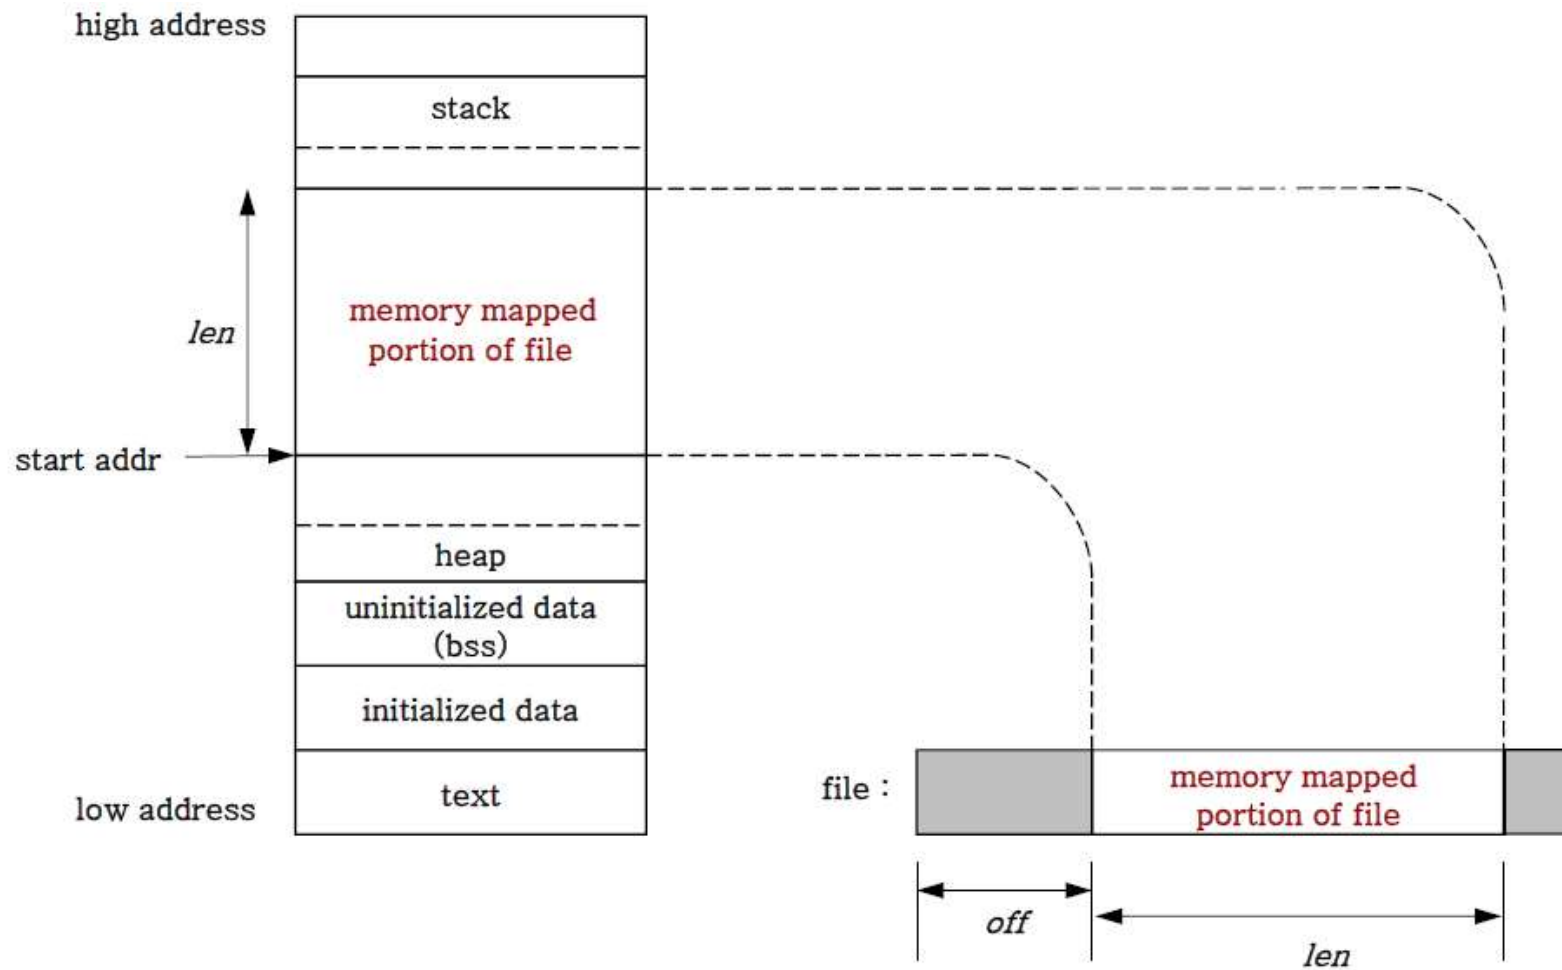
```
#include <sys/types.h>
#include <sys/mman.h>

caddr_t mmap (caddr_t addr, size_t len, int prot, int flags,
int fd, off_t offset);
```

- parameters:
 - *addr* : starting address of a memory area to be mapped to a file
 - *len* : mapping size
 - *prot* : access permission to the memory area
 - *flags* : the mapping scheme
 - *fd* : the file descriptor to be mapped
 - *offset* : the offset in a file where the mapping starts
- return:
 - the real starting address of a memory area to be mapped
 - the kernel may map a file into a different area
 - error : MAP_FAILED



File-memory Mapping by *mmap()*



prot and flags arguments

prot	meaning
PROT_READ	Reading is possible
PROT_WRITE	Writing is possible
PROT_EXEC	Execution is possible
PROT_NONE	Nothing is allowed
* the protection specified for a region has to match the <code>open</code> mode of the file.	

flags	meaning
MAP_SHARED*	Writing to memory is synchronized to the file
MAP_PRIVATE*	Writing is occurred only in another file copy, and the original file is not modified
MAP_FIXED	Resulting map address must be same as the address given as an argument. If not <code>MAP_FIXED</code> , the address argument is just for reference.
MAP_NORESERVE	Do not preserve a swap space
* either <code>MAP_SHARED</code> or <code>MAP_PRIVATE</code> should be specified.	



Release a memory mapping

```
#include <sys/types.h>
#include <sys/mman.h>

int munmap (caddr_t addr, size_t len);
```

- parameters:
 - *addr* : the starting address of a memory area to be unmapped
 - *len* : length
- return:
 - 0 if OK
 - -1 on an error



File copy using mmap file (1)

mmcp.c

```
#include <sys/mman.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <stdio.h>
#include <stdlib.h>
#include <memory.h>
#include <fcntl.h>

#define FILE_MODE (S_IRUSR | S_IWUSR)

void mperr(char *call, int val)
{
    perror( call);
    exit( val);
}

int main( int argc, char *argv[])
{
    int fdin, fdout;
    caddr_t src, dst;
    struct stat statbuf;
```



File copy using mmap file (2)

```
if (argc != 3)
    mperr ("usage : a.out <fromfile> <tofile>", 1);

if ((fdin = open ( argv[1], O_RDONLY)) < 0) {
    fprintf(stderr, "cannot open %s for writing", argv[1]);
    exit(2);
}
if ((fdout = open(argv[2], O_RDWR|O_CREAT|O_TRUNC, FILE_MODE)) < 0) {
    fprintf(stderr, "cannot create %s for writing", argv[2]);
    exit(3);
}

if (fstat(fdin, &statbuf ) < 0)
    mperr("fstat error", 4);

if (lseek(fdout, statbuf.st_size - 1, SEEK_SET) == -1)
    mperr("lseek error", 5);

if (write(fdout, "", 1) != 1)
    mperr("write error", 6);
```



File copy using mmap file (4)

```
if ((src = mmap(0, statbuf.st_size, PROT_READ, MAP_SHARED, fdin, 0)) == MAP_FAILED)
```

```
    mperr("mmap error for input", 7);
```

```
if ((dst = mmap(0, statbuf.st_size, PROT_WRITE, MAP_SHARED, fdout, 0)) == MAP_FAILED)
```

```
    mperr("mmap error for output", 8);
```

```
memcpy(dst, src, statbuf.st_size);
```

```
if(munmap(src, statbuf.st_size) != 0)
```

```
    mperr("munmap(src) error", 9);
```

```
if(munmap(dst, statbuf.st_size) != 0)
```

```
    mperr("munmap(src) error", 10);
```

```
exit(0);
```

```
}
```



Run & Result

```
$ cat myfile
This is a test data.
$ ./a.out myfile mycopy
$ ls -ld myfile mycopy
-rw-rw-r-- 1 shlim shlim 512 Jul 9 22 : 43 myfile
-rw-rw-r-- 1 shlim shlim 512 Jul 9 24 : 13 mycopy
$ cat mycopy
This is a test data.
$
```



Disk Synchronization of a MM Area

```
#include <sys/mman.h>
#include <unistd.h>

int msync (const void *addr, size_t len, int flags);
```

- parameters:
 - *addr* : starting address
 - *len* : length
 - *flags* : control **msync()** operation
- return:
 - 0 if OK
 - -1 on an error



flags argument

flags	meaning
MS_ASYNC	Asynchronous sync request. Do not wait the completion of the sync.
MS_SYNC	Synchronous sync request. Be blocked(waits) until the sync has been completed
MS_INVALIDATE	Invalidate other memory mapping on the same file



msync ex

msync-ex.c

```
#include <stdio.h>
#include <string.h>
#include <errno.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdlib.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/mman.h>

#define MEM_SIZE 64

int main(int argc, char **argv)
{
    int fd;
    char *memPtr = NULL;
    struct stat sb;
    int flag = PROT_WRITE | PROT_READ;

    if (argc != 2) {
        fprintf(stderr, "Usage: %s memPtr\n", argv[0]);
        exit(1);
    }
}
```



msync ex

msync-ex.c

```
if ((fd = open(argv[1], O_RDWR | O_CREAT)) < 0) {
    perror("File Open Error");
    exit(1);
}

if (fstat(fd, &sb) < 0) {
    perror("fstat error");
    exit(1);
}

memPtr = (char *)mmap(0, MEM_SIZE, flag, MAP_SHARED, fd, 0);
if (memPtr == (void *)-1) {
    perror("mmap() error");
    close(fd);
    exit(1);
}

printf("mem(%p),value(%s)\n", memPtr, memPtr);
```



msync ex

msync-ex.c

```
// mem <--> file (synchronization)
while(1){
    scanf("%s", memPtr);
    if (!strcmp(memPtr, "exit")) break;
    if (msync(memPtr, MEM_SIZE, MS_SYNC) == -1)
    {
        printf("mync() error(%s)\n", strerror(errno));
        break;
    }
}

if (munmap(memPtr, MEM_SIZE) == -1) {
    printf("munmap() error(%s)\n", strerror(errno));
}
close(fd);
}
```

