

System Programming

10. POSIX Semaphores

Seung-Ho Lim

Dept. of Computer & Electronic Systems Eng.



Semaphore

- System calls designed by Dijkstra in 1960's
- Used to control the critical section (i.e. resource sharing) b/w multiple threads(processes)
- Basic features
 - A **block/wakeup algorithm** for mutual exclusion.
 - If a process cannot enter a critical section, the process will be **blocked**. (release CPU)
 - When the owner process exits a critical section, it **wakes the waiting process up**. (make it to be **ready**)
 - No waste of CPU time.
 - A **waiting queue** is necessary if several processes are waiting for the permissions for entering their critical sections.



Counting Semaphore

class Semaphore

```
{ private:  
    int value;  
    PCB *queue;  
}
```

Semaphore::Semaphore (n)

```
{ value = n; }
```

Semaphore::wait()

// original name: "P" operation

```
{  
    value--;  
    if ( value < 0 ) {  
        block the calling process;  
        add it to the wait queue of this  
        semaphore;  
    }  
}
```

Semaphore::signal()

// original name: "V" operation

```
{  
    value++;  
    if ( value <= 0 ) {  
        remove the first process from the  
        wait queue;  
        add it to the scheduling queue;  
    }  
}
```



Binary Semaphore

class Semaphore

```
{ private:  
    int value;  
    PCB *queue;  
}
```

Semaphore::Semaphore (1)

```
{ value = 1; }
```

Semaphore::wait()

// original name: “P” operation

```
{  
    if (value == 1) value = 0;  
    else {  
        block the calling process;  
        add it to the wait queue of this  
        semaphore;  
    }  
}
```

Semaphore::signal()

// original name: “V” operation

```
{  
    if ( queue is not empty) {  
        remove the first process from the wait  
        queue;  
        add it to the scheduling queue;  
    } else  
        value = 1;  
}
```

Binary semaphore is a special
case of Counting semaphore!



Mutual Exclusion by Semaphore

- Mutual exclusion using **a binary or a counting semaphore**

Semaphore S(1); // global variable for mutual exclusion, initialized with 1

process/thread A

S.wait ();

critical section;

S.signal ();

remainder section;

process/thread B

S.wait ();

critical section;

S.signal ();

remainder section;



Resource Allocation by a Counting Semaphore

Semaphore Printer(3); // in case that the system has 3 identical printers.


```

:
-----
Printer.wait();
-----
    use a printer; // for a mutually exclusive use of a printer
-----
Printer.signal();
-----
    remainder section;
```



Synchronization with a Semaphore

Semaphore Sync(0); // global var. with the initial value of 0



- In the case that the P_1 reaches the wait-point earlier than P_0 's signal point
 - P_1 will be blocked until P_0 reaches the signal point because the semaphore value is 0.
- In the case that the P_0 reaches the signal-point earlier than P_1 's wait-point
 - P_0 will continue without waiting because the semaphore value is already 1.
- This means that P_0 can proceed after P_1 's sync-point.



Producer/Consumer with a Counting Semaphore

// global variables

Semaphore *mutex(1), num_buffer(n), num_data(0);*

int rear = front = -1; count = 0;

```
while (1) // producer
{
    produce an item;
    num_buffer.wait (); // for sync
    mutex.wait();
    buffer[rear] = pdata;
    rear = (rear + 1) % n ;
    mutex.signal();
    num_data.signal(); // for sync
}
```

```
while(1) // consumer
{
    num_data.wait(); // for sync
    mutex.wait();
    cdata = buffer[front] ;
    front = (front + 1) % n ;
    mutex.signal();
    num_buffer.signal (); // for sync
    process an item;
}
```

1. “Count” variable is unnecessary because we use a counting semaphores !
Counting semaphore **num_data** itself has the number of data.
2. For a single producer and a single consumer, **mutex.wait()** and **mutex.signal()** is unnecessary.



POSIX SEMAPHORES

- NAME

sem_init, sem_wait, sem_trywait, sem_post, sem_getvalue, sem_destroy - operations on semaphores

- SYNOPSIS

#include <semaphore.h>

*int sem_init(sem_t *sem, int pshared, unsigned int value);*

*int sem_wait(sem_t * sem);*

*int sem_trywait(sem_t *sem);*

*int sem_post(sem_t * sem);*

*int sem_getvalue(sem_t * sem, int * sval);*

*int sem_destroy(sem_t * sem);*

- *must be linked with **pthread** library.*



Semaphore initialization

```
#include <semaphore.h>
```

```
int sem_init(sem_t *sem, int pshared, unsigned int value);
```

- parameters:

- *sem* : initializes the semaphore object pointed to by sem.
- *pshared* :
 - 0 means the semaphore is local to the current process (when it is zero)
 - non-zero means that it is shared between several processes (should be located in a region of shared memory b/w the processes)
- *value* : initial value for the semaphore



Semaphore wait operations

```
#include <semaphore.h>

int sem_wait(sem_t * sem);
int sem_trywait(sem_t *sem);
```

■ sem_wait()

- suspends the calling thread until the semaphore pointed to by sem has non-zero count.
- It then atomically **decreases** the semaphore count.

■ sem_trywait()

- is a **non-blocking** version of sem_wait.
- If the semaphore pointed to by sem has non-zero count, the count is atomically decreased and sem_trywait immediately returns 0.
- If the semaphore count is zero, sem_trywait immediately returns with error EAGAIN.



Semaphore post(signal) operations

```
#include <semaphore.h>

int sem_post(sem_t * sem);
int sem_getvalue(sem_t * sem, int * sval);
```

- In POSIX, semaphore signal operation is named ***post***
- **sem_post()**
 - atomically **increases** the count of the semaphore pointed to by *sem*.
 - this function never blocks and can safely be used in asynchronous signal handlers.
- **sem_getvalue()**
 - stores the current count of the semaphore *sem* in the location pointed to by **sval**



Semaphore destruction

```
#include <semaphore.h>

int sem_destroy(sem_t * sem);
```

- destroys a semaphore object pointed by *sem*
 - frees the resources it might hold.
 - no threads should be waiting on the semaphore at the time `sem_destroy` is called.
 - destroying a semaphore that other processes or threads are currently blocked on produces undefined behavior

Example : Mutual Exclusion using Semaphore (1)

semaphore.c

```
#include <semaphore.h>
int cnt=0;
static sem_t hsem;

int main(int argc, char *argv[])
{
    pthread_t thread1;
    pthread_t thread2;

    if (sem_init(&hsem, 0, 1) < 0){
        fprintf(stderr, "Semaphore Initilization Error\n");
        return 1;
    }
    pthread_create(&thread1, NULL, Thread1, NULL);
    pthread_create(&thread2, NULL, Thread2, NULL);
    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);
    printf("%d\n", cnt);
    sem_destroy(&hsem);

    return 0;
}
```



Example : Mutual Exclusion using Semaphore (2)

semaphore.c

```
void *Thread1(void *arg)
{
    int i, tmp;

    for(i=0; i<1000; i++){
        sem_wait(&hsem);
        tmp=cnt;
        usleep(1000);
        cnt=tmp+1;
        sem_post(&hsem);
    }
    printf("Thread1 End\n");

    return NULL;
}
```

```
void *Thread2(void *arg)
{
    int i, tmp;

    for(i=0; i<1000; i++){
        sem_wait(&hsem);
        tmp=cnt;
        usleep(1000);
        cnt=tmp+1;
        sem_post(&hsem);
    }
    printf("Thread2 End\n");

    return NULL;
}
```



Example: Producer & Consumer (1)

prod-cons.c

```
.....
#define MAX_BSIZE  10
int cnt=0;
static sem_t hsem, num_buff, num_data;

void *Producer(void *arg) // producer
{  int i, tmp;

    for(;;) {
        sem_wait(&num_buff);
        sem_wait(&hsem);
        cnt++;
        printf("prod cnt: %d \n", cnt);
        sleep(1);
        sem_post(&hsem);
        sem_post(&num_data);
    }
    printf("Producer Ends\n");
    return NULL;
}
```



Example: Producer & Consumer (2)

```
void *Consumer(void *arg) // consumer
{
    int i, tmp;

    for(;;) {
        sem_wait(&num_data);
        sem_wait(&hsem);
        cnt--;
        printf("cons cnt: %d \n", cnt);
        sleep(1);
        sem_post(&hsem);
        sem_post(&num_buff);
    }
    printf("Consumer Ends\n");
    return NULL;
}
```



Example: Producer & Consumer (3)

```
int main(int argc, char *argv[])
{
    pthread_t thread1;
    pthread_t thread2;

    if (sem_init(&hsem, 0, 1) < 0){
        fprintf(stderr, "Semaphore Initialization Error\n");
        return 1;
    }
    if (sem_init(&num_buff, 0, MAX_BSIZE) < 0){
        fprintf(stderr, "Semaphore Initialization Error\n");
        return 1;
    }
    if (sem_init(&num_data, 0, 0) < 0){
        fprintf(stderr, "Semaphore Initialization Error\n");
        return 1;
    }
    pthread_create(&thread1, NULL, Producer, NULL);
    pthread_create(&thread2, NULL, Consumer, NULL);
    pthread_join(thread1, NULL);
    pthread_join(thread2, NULL);
    printf("%d\n", cnt);
    sem_destroy(&hsem);

    return 0;
}
```

